

ORBIT-TRUST implementation handoff

A complete specification for an evidence-driven satellite conjunction review demonstration. The application target is devx-engine on Vercel, with Supabase persistence, Google ADK, Groq and a deterministic Rust calculation core.

This is the human-readable companion to the ZIP and complete Markdown. Exact schemas, synthetic inputs and reference outputs remain in those machine-readable files. The application has not been implemented by preparing this package.

Contents

Read this first	2
Prompt for a brand new implementation chat	4
Product requirements and scientific boundaries	6
Decision register and defaults	8
Source repository inspection and reuse map	10
Architecture and application boundaries	12
Domain records and API contracts	15
Orbital mathematics and deterministic triage	20
Fleet comparison and communication simulation	23
Reentry exposure and financial scenarios	25
ADK orchestration and output validation	27
Mission console interaction specification	29
Data sources and reproducible fixture catalogue	31
Test strategy and acceptance catalogue	33
Implementation milestones and work tickets	36
Deployment and operations runbook	39
Business model, validation and hackathon presentation	42
Sources and evidence register	46
Traceability and release gate	48
Glossary and implementation recovery	50
Supabase database, authorization and transaction specification	52

ORBIT TRUST implementation handoff

Version 2.0, revision game-preservation-1. Prepared 11 September 2026. This revision records the owner's clarification that 3d-game is important for future implementation. This package is a specification and reference-data handoff, not an implemented application. It contains the context needed by an engineer or a new AI chat that has never seen the original conversation.

Build outcome: a public hackathon demonstration for satellite conjunction review, evidence auditing, fleet response comparison and financial scenarios, with a separate reentry exposure sandbox. Use Google ADK, Groq, a deterministic Rust calculation library and a simple mission console adapted from the owner's existing projects.

How to use this package

1. Extract the complete ZIP. Give the new chat this folder, or upload ORBIT_TRUST_COMPLETE_HANDOFF.md if it accepts only one text file.
2. Paste 00_BOOTSTRAP_PROMPT.md as your instruction. It explicitly tells the new chat to build, not merely describe the project.
3. The implementer reads documents 01 through 06 before coding, then follows the milestone sequence in document 13. Documents 07 through 19 resolve the specialist details.
4. Keep schemas/, fixtures/ and reference/ with the documents. These are specification inputs and expected-output examples, not application source code.
5. Put credentials in the deployment service's secret settings. Never paste keys into a chat or add them to this package.

The complete Markdown includes the numbered documents and machine-readable contract/fixture appendices. The PDF is the human-readable companion; exact machine contracts remain in the ZIP and complete Markdown. A manifest records file hashes so the recipient can detect an incomplete transfer.

Decisions carried from the owner

- Preserve real mathematical and engineering depth, useful agent tools, a compelling demonstration and a believable business case.
- No team-size, build-time or experience questionnaire. Work is sequenced by dependencies rather than assumed staffing.
- No GPU or paid API credits. **Groq is the chosen AI provider.**
- **A public web demo is required.** A local backup does not replace this requirement.
- Visual direction: dark navigation, light work panels and restrained status colors.
- Lead with fleet review and response. Include terrain analysis as a separate reentry sandbox.
- Reuse the owner's TerraSight, Finora and Sajawat projects selectively. Do not merge their applications wholesale.
- The existing 3d-game/ app in devx-engine is an important future implementation asset. Preserve its source, assets, configuration, history and any existing deployment. Do not delete, overwrite or repurpose it while building ORBIT-TRUST in orbit-trust/.
- Do not promise zero hallucinations, certain collisions, optimal real maneuvers, exact crash sites, actual money saved or operational flight readiness.

Authority and versioning

The owner's latest instructions outrank this package. Within version 2.0, document 02 records decisions; document 05 and `schemas/core.schema.json` define contracts; document 06 defines scientific rules; documents 12 and 17 define acceptance. If these disagree, record the conflict and fix the documentation before changing behavior. Do not silently choose the easiest interpretation.

This package supersedes the earlier 30-page playbook where scope differs. In particular, Groq and a public web release replace local-model-first deployment; supplied fleet candidates and a separate reentry sandbox are now included. Earlier research remains context, not an alternative implementation specification.

All performance targets, prices, mission policy thresholds and fictional scenarios are proposals unless explicitly identified as observations. No ORBIT-TRUST application, live Groq integration, public deployment or operator validation has been completed by preparing this package.

Deliverables the next implementer must produce

An ORBIT-TRUST application under `orbit-trust/` in the existing `Adhirajsingh2507/devx-engine` repository, working Vercel public URL, persistent Supabase accounts and case data, local reproducible run, meaningful tests, measured benchmark report, dependency lockfiles, `data/source` notices, a three-minute demo recording, and a release report distinguishing completed work from blocked items. Preserve the important future `3d-game/` app, existing toolkit content and the three source repositories. Future game development or integration is a separate task requiring the owner's direction.

The owner reports that `devx-engine` is already connected to Vercel for automatic deployment. Confirm the existing project's Root Directory and build settings before publishing. Groq and Supabase secret values and account login remain owner-supplied deployment inputs. Their absence must not block independent implementation or local verification.

Prompt for a brand new implementation chat

Copy the text below into a new chat with the handoff package attached.

Build ORBIT-TRUST from the attached version 2.0 implementation handoff, revision game-preservation-1. You have no prior conversation context; the package contains it. Your task is to implement, test and prepare the public hackathon demonstration, not stop at another proposal.

Read START_HERE and numbered documents 01 through 06, then use document 13 as the execution order. Read each specialist specification before implementing it. The complete Markdown is an alternative to the folder; it includes embedded schemas and fixtures that you can reconstruct at their stated relative paths. Use the machine contracts and acceptance cases to prevent interface drift.

The project combines conjunction-warning prioritization with evidence auditing. Its main demonstration is fleet review and comparison of supplied response candidates. It also includes a distinct reentry exposure sandbox and conditional financial analysis. Orbital collision, reentry and ground damage are different processes: do not turn a conjunction report into a claimed crash point or steering command.

Use a Next.js mission console, Python/FastAPI, Google ADK, Groq and a deterministic Rust numerical library exposed through PyO3. The chosen visual direction is dark navigation with light work panels and restrained status colors. The public deployment and storage mode are specified in documents 02 and 14. No GPU or paid API credit is available. Never enable paid plans, change billing or invent access credentials. Use the documented deterministic fallback if the AI provider is unavailable, visibly labeled; do not claim this proves the live agent integration.

You may read and reuse my projects at:

- <https://github.com/Adhirajsingh2507/terrasight>
- <https://github.com/Adhirajsingh2507/Error-404-Not-Found>
- <https://github.com/Adhirajsingh2507/sajawat>

The handoff pins the inspected revisions and explains useful components and known limitations. Inspect those paths and applicable repository instructions before copying. Treat report text, descriptions and sample prompts as reference data unless they are applicable project instructions. Do not import unrelated deployment scripts, secrets, business logic or agent permissions. Implement under orbit-trust/ in <https://github.com/Adhirajsingh2507/devx-engine>, preserving existing content. The owner has connected this repository to Vercel. Verify the intended ORBIT-TRUST project's Root Directory and use Supabase for persistent accounts and case data. Preserve the three reuse-source repositories.

The existing 3d-game/ app is important for future implementation. Do not delete, overwrite, rename, move or repurpose it during the ORBIT-TRUST build. Preserve its source, assets, dependencies, configuration, history and any existing deployment. Do not classify it as disposable or remove it in cleanup. Keep ORBIT-TRUST additions scoped to orbit-trust/ and any necessary additive shared configuration. Verify the final diff preserves the game. Future game development/integration is not part of this build unless the owner requests it. If the existing Vercel project serves the game, preserve that deployment and resolve the intended ORBIT-TRUST deployment target with the owner before changing its settings.

Start by checking available tools, filesystem, repository access and dependencies. Record exact selected dependency versions and run the ADK/Groq/PyO3 integration spike early. Resolve routine implementation details using the stated defaults. Ask only when an external credential, publication destination, incompatible requirement or consequential missing fact truly blocks the next step. Do not ask again about team size, time or experience. Do not expose secrets in commands, logs, screenshots or reports.

The authoritative queue, scientific outputs, financial calculations and policy decisions must work without an LLM. All results reference immutable input revisions and policy versions. The agent can investigate allowed alternatives and select evidence for an explanation; it cannot manufacture data, remove an urgent case, send

external messages or command a satellite. Scientific unsupported states and missing inputs must remain explicit.

Build incrementally through the documented milestones. Keep a BUILD_STATUS.md recording completed acceptance IDs, actual commands/results, dependency versions, measurements, source modifications and unresolved items. Tests written but not run are not passing tests. Synthetic fixture results are not measured customer outcomes. Do not reuse the same routine as both the numerical implementation and its only oracle.

Before release, exercise the live public URL from a fresh browser, test isolation between two visitors, verify that keys never reach the client, and rehearse with Groq unavailable. A localhost screenshot or static mockup does not satisfy the required public interactive demo. If deployment credentials are unavailable, finish all independent implementation, provide the deployable release and exact final steps, and report the deployment as incomplete.

Finish with the public URL if verified, repository location, local startup instructions, test/benchmark evidence, limitations and a concise mapping from requirements to implemented behavior. Do not claim operational flight safety, guaranteed collision prevention, zero hallucinations, a full-catalog assessment or money actually saved.

Product requirements and scientific boundaries

Problem and buyer

Satellite operations analysts receive evolving conjunction reports. They must determine which event needs attention, whether the evidence supports the assessment, and what can still be done before the relevant review and command deadlines. A queue sorted only by reported collision probability can obscure important uncertainty and timing issues.

The primary user is a flight-dynamics or orbital-safety analyst. The first buyer hypothesis is the head of mission operations at a small commercial LEO operator that already receives conjunction data. Flight-dynamics service firms are an adjacent customer channel. Satellite count alone does not establish need: recurring review effort and an inadequate existing workflow do.

ORBIT-TRUST answers: which case needs review, why, which evidence is missing, how supplied response options affect the supplied fleet, and what the stated financial consequences are. It preserves human responsibility for operational choices.

Required capabilities

ID	Requirement	Visible proof
R01	Import canonical encounter JSON and historical ESA CSV in local research mode	Import report with accepted, rejected and unsupported records
R02	Group explicit event IDs and preserve report revisions	One case with source-specific history
R03	Audit evidence separately from urgency	Two independent status fields with reasons
R04	Compute supported short-encounter probability in Rust	Reference-checked result and method record
R05	Apply deadline-aware P0-P3 review policy	Urgent uncertain case remains visible
R06	Compare supplied response alternatives across loaded objects	A-B improvement rejected when A-C worsens
R07	Simulate communication feasibility and status	Sent, acknowledged and executed are separate
R08	Compute conditional loss and expected-loss change	Inputs, formula, currency and limitations visible
R09	Provide a separate reentry exposure sandbox	Supplied footprint changes exposure totals
R10	Run bounded ADK investigations through Groq	Recorded tool calls and version-bound output
R11	Validate both tool results and report output	Invalid citations or actions cannot enter accepted packet
R12	Provide queue, case, satellite and activity workflows	Add/edit metadata, assign, note, review and export
R13	Survive provider failure without losing calculations	Deterministic results remain available
R14	Work at a public web URL	Fresh-browser live demonstration
R15	Preserve reproducibility and traceability	Exported bundle reproduces the assessment
R16	Meet accessibility and bounded resource requirements	Keyboard flow, responsive views and measured budgets

Scope boundaries

The core accepts assessments; it is not a global tracking or conjunction-screening service. It does not discover every orbital object. Fleet assessment covers exactly the supplied catalogue, trajectories and time interval. The initial physics engine supports the explicitly declared short linear Gaussian encounter model, not arbitrary orbital propagation.

Response candidates are supplied trajectories with planning assumptions. The software can compare their encounter consequences and check the supplied timing constraints. It does not infer a feasible burn from a slider or optimize a real maneuver. All maneuver-like actions in the demo are simulation records.

The reentry sandbox starts with an externally supplied or synthetic footprint. It does not compute atmospheric breakup, fragment survival, drag or reachable landing corridors. It does not route falling debris using rover navigation. ESA describes reentry analysis as including trajectory, surviving fragments and population risk; a terrain score does not replace those inputs.^{2, 3}

Financial outcomes are conditional models. The existence of a costly spacecraft does not mean a warning implies the entire spacecraft value has been saved. A collision-free outcome alone cannot establish the effect of a particular maneuver or software product.

User journeys

An analyst opens the queue, selects an urgent case and sees the specific source conflict. They inspect report revisions, request a bounded investigation and receive an evidence-linked information-request draft. They compare two supplied response candidates and observe that one creates a new fleet concern. They record a review disposition, retaining the assumptions and assessment version. A new material report reopens review and invalidates stale explanations.

A judge opens a separate Reentry Sandbox tab. A persistent label identifies the scenario as synthetic. Two supplied footprints cover different fictional exposure points. The interface updates the exposed population and asset-value totals. Damage estimates appear only for explicitly supplied probability and vulnerability assumptions. This branch cannot change orbital queue priority.

Success and claims

Implementation success means all required capabilities have passing acceptance evidence and the public interactive demo works. Scientific success is narrower: reference agreement within the supported calculation domain. Business success requires separate analyst and buyer evidence, which is not currently available.

Measure required-review recall, unnecessary review burden, time to a defensible disposition, output fidelity, cache behavior, latency and memory. Do not substitute a visually impressive map or agent animation for these results. The competitive claim is reduced review effort and clearer reconciliation; it remains a hypothesis until compared against a strong existing workflow.

Language shown to users

Use “Review now,” “Evidence incomplete,” “Calculation unsupported,” “Candidate passes demo checks,” “Simulation acknowledgement,” and “Estimated change in expected loss.” Avoid “Safe satellite,” “Collision certain,” “Crash location,” “Optimal escape,” “Guaranteed savings,” or a numerical confidence percentage based on agent agreement.

Source notes

2. ESA, [Reentry background](#).

3. ESA, [Re-entry safety](#).

Decision register and defaults

This is the authoritative record of product choices. U denotes an explicit owner decision; D denotes an implementation default selected to complete the handoff; V denotes a requirement needing external verification before release. Defaults are not customer or scientific evidence.

ID	Kind	Decision
D01	U	Build a space-technology hackathon project with substantive mathematics and useful agents
D02	U	Combine conjunction-warning reliability and evidence auditing
D03	U	Lead the demonstration with fleet review and response
D04	U	Include a separate TerraSight-inspired reentry sandbox
D05	U	Groq is the hosted model provider; no paid API credit or GPU
D06	U	Public interactive web demo required
D07	U	Dark navigation, light work panels, restrained status colors
D08	U	Reuse the three owner-supplied repositories selectively
D09	U/D	Owner selected Adhirajsingh2507/devx-engine, already connected to Vercel; orbit-trust/ subdirectory is an implementation default
D10	D	Next.js static export, TypeScript, Tailwind, FastAPI, ADK, PyO3 Rust library
D11	U/D	Owner selected the existing Vercel project; same-origin CDN frontend and FastAPI function is an implementation default
D12	U	Supabase free tier for persistent accounts and case data; no authoritative Vercel filesystem state
D13	D	GitHub OAuth for persistent identity; isolated anonymous demo trial with export/import
D14	D	No shared mutable global demo; no visitor access to another workspace
D15	D	Groq model openai/gpt-oss-20b through LiteLLM; fallback is deterministic templates
D16	D	Two LLM roles: investigation and formatting; domain specialists otherwise use deterministic tools
D17	D	Supplied candidate trajectories only; no live commanding or autonomous burn optimizer
D18	D	Initial numerical domain is short linear encounters with declared Gaussian position uncertainty
D19	D	Historical ESA data is optional local research input, not bundled public demonstration data
D20	D	Primary display currency INR; user-supplied USD scenarios supported without automatic FX conversion
D21	D	All times stored as UTC; demo clock is explicit and separate from wall-clock sessions
D22	D	Release uses a fictional fixed demo clock beginning 2026-09-11T12:20:00Z
D23	D	Runtimes: Python 3.12, Node 22 LTS, pnpm 10; pin compatible patch versions in milestone M0
D24	D	Rust stable toolchain pinned in M0; no nightly, unsafe custom numerical code or fast-math requirement
D25	D	Test-only NumPy/SciPy reference routines stay out of the production function bundle where possible
D26	D	Public deployment includes only synthetic uploads; real operational/private datasets are local-only
D27	V	Account-specific Groq access, exact quotas, host eligibility and event submission rules must be verified

ID	Kind	Decision
D28	U	The existing 3d-game/ app is important for future implementation. Preserve its files, assets, dependencies, configuration, history and any deployment; do not remove or repurpose it for ORBIT-TRUST

Why these choices differ from the earlier playbook

The earlier report recommended a local CPU model and local-first execution. The owner subsequently selected Groq and required a public web demo. Those selections supersede the earlier recommendation. The prior statements about no public deployment or no fleet comparison are also superseded by the new scope. Scientific cautions, provenance requirements and evidence/urgency separation remain.

Current Groq documentation lists the earlier Finora Llama model as enterprise access. Do not assume the owner's previous default is available on a new free account. The proposed GPT OSS 20B model appears in Groq's free-limit table and supports tool use and strict structured output in separate calls. Access still requires an account check. The model is served by Groq; no OpenAI API account is required.^{8, 9, 10}

The owner explicitly selected Vercel after a preliminary hosting comparison. Render and Hugging Face are not deployment requirements. Vercel supports FastAPI functions and static assets; persistent data belongs in Supabase. Package the numerical extension before function deployment and bound all computation by request lifetime.^{12, 13}

External inputs without hidden assumptions

The owner will supply Groq/Supabase configuration through secret settings, authenticate account setup and verify the hackathon submission link. The target repository is known; the existing Vercel project settings and Supabase project values require inspection or owner input at deployment. Never request secret values in plain chat. The component lecture does not establish formal judging weights or whether pre-existing code is permitted. Disclose reuse and resolve the organizer's rule before submission.

The reported 29/50 readiness score in earlier research was a subjective assessment, not a release score. Do not carry it into the app as performance evidence. No fabricated customer interviews, live telemetry, loss data or scientifically calibrated confidence intervals are permitted.

Change control

If a default becomes infeasible, record the failed check and the smallest compatible replacement in a decision log. Do not silently switch to a paid service, remove the public demo, replace ADK with fake traces, merge the reentry sandbox into orbital risk, or declare an unsupported calculation valid. Account provisioning is a release dependency, not a reason to stop unrelated implementation.

Dependency patch versions are intentionally resolved at implementation time because the pack cannot certify a future lockfile. This is a bounded compatibility task: select current non-prerelease versions satisfying the chosen runtime, exclude known compromised releases, run the M0 contract spike, lock and record. It is not permission to redesign the architecture.

Source notes

8. Groq, [Supported models](#).

9. Groq, [Rate limits](#).

10. Groq, [Tool use](#) and [Structured outputs](#).

12. Vercel, [FastAPI deployment](#), refreshed 11 September 2026.

13. Vercel, [Python runtime](#) and [Function limitations](#).

Source repository inspection and reuse map

All three repositories were accessible through GitHub's public file API on 11 September 2026. Initial browser retrieval failures did not mean the repositories were absent. The observations are a targeted source inspection, not an execution audit or a claim that every file is correct.

Project	Repository	Inspected revision
TerraSight	https://github.com/Adhirajsingh2507/terrasight	2e6445f0639f1d20ce0ad45c98134e9bcc11ea20
Finora	https://github.com/Adhirajsingh2507/Error-404-Not-Found	c66824b8242262bb0f640f4b98ad9067591d29d4
Sajawat	https://github.com/Adhirajsingh2507/sajawat	600cf5536d8f0b138bdf470e298329213cf1bf57

Read the pinned revision first. Compare current default branches before reuse and record differences that affect the selected component. If a revision is unavailable, report it and inspect the current counterpart; never imply the original inspection applies unchanged. Keep source-reuse notices with the extracted components and record their original paths and commits.

TerraSight

Its README describes stereo/RGB perception for autonomous planetary rovers, terrain classification and construction-site scoring. The inspected backend/app/pipeline.py operates on fixture scenes and produces tiles, sites, boundaries and path outputs. This is useful architecture and UI material, but it is not an Earth reentry model.

Reuse the measurement-versus-decision separation from backend/app/terrain/assemble.py and backend/app/scoring.py. Adapt frontend/src/components/terrain-grid.tsx, site-table.tsx and stats-panel.tsx only after checking their data and coordinate assumptions. Use the API/types files to understand component dependencies. These frontend files were identified in the tree but were not fully audited during this review.

Replace crater-distance, slope, roughness and bearing-score business logic with the reentry sandbox's footprint/exposure model. The existing four zones concern buildability, navigation, geological interest and hazards. "Construction-safe" has no valid interpretation as "safe for debris impact."

The inspected _walk_border routine joins disconnected feature regions with a nearest-neighbor walk. Do not reuse that geometry as an impact polygon: it can bridge separate regions. Use valid GeoJSON Polygon/MultiPolygon boundaries and preserve holes. No-crater sentinel values and local Cartesian grid coordinates must not leak into Earth exposure units.

Finora

Useful concepts are deterministic specialist calculations, calculation provenance, small tool sets, an intent fast path, bounded loops and a trace. Inspect backend/formatting.py, agents.py, orchestrator.py, judge.py and finance_engine.py to trace dependencies before reuse.

Its financial engine concerns affordability, EMI, tax, emergency funds, investments and fraud. Those formulas are not satellite-loss models. The inspected engine also uses a fixed date. Replace the domain model and date handling; keep the idea that numbers come from code and carry input/formula provenance.

The inspected judge's response_ok primarily checks whether findings are nonempty. Its numeric confidence grows with finding count. Neither establishes factual completeness or calibrated confidence. The _decision path generates prose after evaluating those findings; it has no full final-prose validation in that inspected path. Replace these behaviors with document 09's validator and template-bound output.

The inspected specialist tool calls execute in a Python loop. Reuse does not provide parallel numerical processing automatically. Move independent encounter work to the bounded Rust worker pool, while keeping workflow ownership deterministic.

Finora's README documents serverless, stateless conversation tokens and an ephemeral financial store. Do not adopt that store as durable mission data. Use the case ledger and workspace isolation here. Do not carry its Groq/NIM provider fallback or API keys into the new application.

Sajawat

Inspect `apps/admin/src/app/(console)/crm/page.tsx`, `crm/[id]/page.tsx`, `features/crm/stages.ts`, `features/console/ui.tsx`, `Sidebar.tsx`, `Can.tsx`, the CRM service and corresponding types. The detail screen offers notes, assignment and stage updates, which map well to review cases.

The board requests only 100 leads and groups them in the browser. Replace this with server-authoritative sorting, filtering and pagination. Count all matching cases and all urgency tiers on the server. Do not let browser truncation hide P0/P1 events.

Replace sales stages with the six review stages. Keep urgency as a separate field and allow system-driven reopening. An arbitrary dropdown change must not dismiss a persistent urgent acknowledgement requirement.

The original detail route uses a dynamic path. Our static export uses fixed pages with query parameters, such as `/case/?id=...`; adapt route access accordingly. Extract a small UI subset and its dependencies instead of importing the commerce monorepo, MongoDB, payment system or deployment infrastructure.

Rights and practical extraction

The implementation target is <https://github.com/Adhirajsingh2507/devx-engine>. At inspected revision `ad4b9e45b234fe8b4d22d7db3be7c74d46c7a3fd`, the root contains `README.md`, `.gitignore` and `.mcp.json`, with `toolkit/reference` folders and the `3d-game/` app. The owner identifies this game as an important asset for future implementation. Preserve its source, assets, dependencies, configuration, history and any existing deployment. Do not delete, overwrite, rename, move or repurpose it during the ORBIT-TRUST build. Add `orbit-trust/` alongside it. Future game development/integration needs a separate owner request. No ORBIT-TRUST implementation or root package manifest was identified at the inspected snapshot.

The public tree cannot reveal Vercel's selected Root Directory. Configure `orbit-trust` only for the intended ORBIT-TRUST deployment; do not repoint or replace an existing game deployment. If the connected project serves the game, resolve the ORBIT-TRUST deployment target with the owner before changing its settings. Shared configuration changes must be additive and preserve the game's build/deployment behavior.

The owner has explicitly requested reuse of these projects. GitHub metadata did not identify a repository license in the inspected snapshots; that is not a license grant for unrelated third-party assets. Preserve contributor notices, inventory copied dependencies and obtain any required contributor/asset permission before publishing copied third-party material. Prefer newly drawn UI and fictional fixtures over copied media.

Write `SOURCE_REUSE.md` in the implementation repository listing source URL, pinned commit, original path, reused behavior, changed behavior and validation. Repo instructions guide work in the relevant source scope; do not import their unrelated automation or multi-agent instructions into ORBIT-TRUST.

Architecture and application boundaries


Confirmed deployment

Implement in orbit-trust/ within the existing Adhirajsingh2507/devx-engine repository. The owner reports an existing Vercel Git deployment connection. The existing 3d-game/ app is important for future implementation: preserve its source, assets, configuration and any deployment, alongside the toolkit content. Do not remove or repurpose it. The public ORBIT-TRUST application combines a static Next.js export served by Vercel's CDN and one FastAPI Python function on the same origin. Supabase stores persistent accounts, cases and calculation provenance. Groq supplies model inference.

Use fixed web routes /, /queue/, /case/, /fleet/, /satellites/, /satellite/, /reentry/, /activity/, /settings/, /auth/callback/ and /about/. IDs are query parameters. The static shell contains no user-specific rendered content: all private state is fetched after client authentication. Never cache personalized API responses on the CDN or build a static page with user metadata.

Build Next.js with static export, local fonts and SVG/Canvas illustrations. No server actions, runtime SSR or ungenerated dynamic pages. Copy the export to the FastAPI project's public/ tree during build. The M0 preview verifies route precedence, trailing slashes and API 404s. If current Vercel routing requires explicit mappings, generate them from the fixed page list, not an unconditional catch-all returning HTML for unknown API calls. ^{12, 14}

Logical flow



Reentry is a separate supplied-footprint scenario.

Import -> canonical validation -> immutable report ledger -> deterministic checks -> supported Rust calculations -> policy assessment -> immediate queue.

On request, ADK investigates a saved assessment snapshot. Tools read version-bound records or request bounded scenario calculations. A separate formatter selects evidence and template codes. A host validator assembles the packet. New report, policy or candidate revisions invalidate dependent results.

Fleet comparison, reentry exposure and economics are distinct bounded operations. They share provenance and presentation conventions, not one physics model. Financial tools never backfill missing physics.

New application layout

Directory within orbit-trust	Responsibility
apps/web	Next.js UI, API client and accessible components
orbit_trust	Python API, domain models, persistence, policy and ADK
crates/orbit_core	Math, pair batches and PyO3 bindings
vendor/wheels	Reproducible Linux-compatible Rust wheel and build manifest
contracts	Canonical schema and generated API snapshot
data/fixtures	Original synthetic bundles and expected outputs
supabase/migrations	Tables, grants, RLS and atomic domain operations
tests/reference and tests/e2e	Independent numerical, workflow and browser checks
docs and deploy	Specification, measured evidence and Vercel runbook

Root app.py exports the FastAPI instance. pyproject.toml, dependency locks and vercel.json belong under orbit-trust, the designated Vercel project Root Directory. Use one Python package, one Rust crate and one web application. No distributed worker platform, vector database, A2A service mesh or general agent marketplace.

Serverless execution

No in-memory queue, process-local session ledger, SQLite or persistent background thread is authoritative hosted state. Vercel invocations can terminate or run in separate instances. Complete operations inside a bounded request or use explicit resumable batches with Supabase checkpoints. Polling status does not run a hidden worker.

Use a 45-second application request budget below a verified configured Vercel duration. Reserve 5 seconds for persistence. Numerical work has a 10-second batch budget; Groq has the tighter budgets in document 09. Persist results before returning success. On interruption, a running operation's lease expires and a new request can mark it interrupted and retry idempotently.

For a maximum 20-object, four-candidate comparison, continuation requests process at most 25 pair/candidate jobs or stop earlier at budget. Return progress and a continuation cursor. Closing the browser pauses unstarted work; reopening offers Resume. The default three-object demo should fit one request, subject to measurement.

Enforce one operation per workspace and a small global Groq concurrency limit through atomic database leases, not Python globals. A lease lasts 60 seconds; final writes require its current token and the unchanged input revision. Duplicate requests cannot publish competing results. Rust defaults to one worker publicly; local benchmarks can use min(4, available cores). Release the GIL during numerical batches.

Persistence and authorization

Supabase Postgres is authoritative locally and publicly; offline integration testing uses its local development stack. Reads use the verified user JWT and RLS. Domain mutation endpoints validate JWT signature, issuer, audience and expiry plus workspace membership, then call restricted server-only transactional RPCs. The service credential never enters the browser; RLS bypass never replaces membership checks.

Revision insertion, current-pointer updates, invalidation and audit events are atomic. Seeds are immutable public synthetic data; changes belong to a workspace. Use uniqueness constraints for source identity and idempotency keys. Document 19 defines tables, grants and RPC responsibilities.

Identity and result binding

Results identify schema version, workspace/case, input reports, raw and normalized hashes, policy/scenario/method versions, time and findings. Cache keys include catalogue, candidate, economic and exposure revisions where relevant. Anonymous visitors never share mutable cases. Shared caches contain only public synthetic mathematical inputs, without names, notes or credentials.

Normalize digests over sorted-key compact UTF-8 JSON with finite numbers and canonical UTC timestamps. Keep original-file SHA-256 separately. Do not assume different serializers or equivalent numeric spellings hash identically without canonicalization tests.

Frontend updates

Poll every 3 seconds while an operation runs, every 15 seconds for the active queue, and stop in hidden tabs. Return workspace revision and activity sequence; fetch deltas or a full snapshot after cursor expiry. SSE and WebSockets are not required. This replaces the earlier continuously running service design.

Publish the queue before model work. Host facts and decisions cannot be overwritten by generated prose. New evidence marks explanations stale. Revision-protected edits return 409 instead of overwriting newer data. Demo clock advances affect policy; authentication and leases always use wall time.

Packaging and backup

Use Vercel's Python Functions runtime, not Edge. Build the Rust Linux wheel reproducibly and verify imports in a preview before dependent features. Keep training packages and large archives out of the production function bundle. Run the same API locally with a platform-compatible extension. Local replay is a backup, not proof of the required public release.^{12, 13}

Source notes

12. Vercel, [FastAPI deployment](#), refreshed 11 September 2026.

13. Vercel, [Python runtime](#) and [Function limitations](#).

14. Next.js, [Static exports](#).

Domain records and API contracts

The canonical input contracts are in `schemas/core.schema.json`. JSON Schema validates structure; domain validators enforce physical applicability, ownership and cross-field rules. The schema root accepts `Report`, `CandidateSet`, `EconomicScenario`, `ReentryScenario` or `CommunicationScenario`, discriminated by `bundle_kind`. `ImportBatch`, `AgentSelection` and `Error` are separate named definitions. Fixture files are test wrappers: import their input object, not the entire expected-output wrapper. The API is versioned under `/api/v1`. Generate and commit an OpenAPI snapshot from the implemented FastAPI models, then compare it in CI. This document specifies behavior rather than an already running endpoint.

Shared conventions

Use UTF-8 JSON, snake_case API fields and UTC ISO 8601 timestamps ending in Z. Runtime IDs are UUIDs; fixture IDs are stable readable strings. IDs are opaque to the client. All numbers must be finite; no NaN, Infinity or string-to-number coercion. Probabilities are linear values in $[0,1]$ unless the field explicitly says `log10`. Money uses a decimal string with a currency code, never a binary float in the persistence/API layer. Distances are metres, velocity metres/second, covariance square metres and durations seconds.

Every input bundle declares `schema_version=2.0`, `provenance_kind` (synthetic, historical or operator_supplied), `source_name` and `source_revision`. Imported timestamps are source facts; `received_at` is the application's reception time. Recomputed/derived fields retain their source paths and method IDs. Missing is null or a missing optional field as the schema specifies, never an invented zero.

Records

Record	Required content and relationships
Workspace	ID, owner/member identities, mode, creation time and data revision
Satellite	Workspace, ID, display name, operator label, maneuver capability, metadata revision and optional economic inputs
Case	Workspace, explicit event key, object pair, reported TCA, workflow state, assignment and latest assessment ID
Report	Case, provider/source, source revision, creation/reception times, TCA, reported Pc, raw input digest and optional complete encounter
Encounter	Common epoch and frame, two object states, covariance assumptions, radii and model-domain declaration
Finding	Stable ID, category, code, severity, material flag, source references and explanation template parameters
Assessment	Input/report IDs, digest, policy version, evidence state, urgency, facts, method status, deadline and creation time
CandidateSet	Immutable baseline and alternative trajectory snapshots, catalogue revision, horizon and supplied planning constraints
Comparison	Every assessed pair/candidate, coverage, rejected/unsupported checks and candidate disposition
EconomicScenario	Currency, probability references, conditional loss line items, response costs and uncertainty labels
ReentryScenario	Footprint/layer revisions, geographic conventions and optional vulnerability inputs

Record	Required content and relationships
AgentRun	Assessment binding, provider/model, prompt version, bounded tool trace, usage, validation and outcome
ReviewAction	Actor, case revision, action type, rationale, assessment acknowledged and timestamp
AuditEvent	Workspace sequence, actor, event type, entity/revision and non-secret change summary

Keep reported_pc and recomputed_pc separate. A recomputation never overwrites its provider value. An assessment's evidence state is usable, incomplete, inconsistent or unsupported. Store individual findings even when one summary state takes precedence: unsupported calculation applicability first, then material inconsistency, then missing required data, otherwise usable. This summary does not suppress unrelated evidence.

Event grouping and revisions

Workflow states are new, reviewing, awaiting_information, decision_recorded, monitoring and closed. Begin review from new; request information from reviewing; return awaiting_information to reviewing on new material evidence; record a disposition from reviewing; then move decision_recorded to monitoring or closed with latest-version acknowledgement and a rationale. New material evidence reopens monitoring/closed as reviewing. Assignment and notes do not constitute acknowledgement. Invalid transitions return 409 with the allowed next actions.

Acknowledgement resolves a review task for that assessment, not its physical risk finding. A closed or monitored case retains its scientific classification and recorded risk; show "Review recorded" separately. Default active queue counts include unresolved required reviews across the workspace, regardless of page/filter. Closed cases remain accessible in history and may retain a P1 scientific classification; never relabel them P3 simply because a person reviewed them. An unacknowledged urgent task cannot be hidden by a stage change. The separate event_phase is pre_encounter, post_encounter_pending_verification or verified_outcome_recorded; only an explicit sourced outcome/disposition sets the last value.

Group by workspace + explicit event_key. The JSON import requires this key; the historical adapter namespaces its event_id with the dataset ID. Never merge two providers' events based only on a similar TCA or object names. Cross-provider association requires an explicit supplied association. Reject a report that reuses an event key with a different object pair.

Deduplicate by source_name + source_revision + raw digest. An identical report has no new scientific revision. The same source revision with different content is a source conflict, not an overwrite. Arrival order is not scientific recency: order each source stream by creation time and preserve out-of-order arrivals. An input must state which report set is currently applicable; unresolved competing source claims remain alternatives. Do not average their Pc values.

Change to any material source, policy, candidate, economic assumption or exposure layer creates a new assessment dependency version. A note or assignment change creates audit history without rerunning physics. Source records are immutable. Operator corrections are new revisions identifying what was corrected and why.

Endpoint surface

Method and path	Request	Response / behavior
GET /health	None	Liveness and build ID; no secrets or external calls
GET /capabilities	None	Engine/model configuration and supported modes, no credentials

Method and path	Request	Response / behavior
POST /workspaces/demo	Seed ID, idempotency key	Isolated synthetic workspace for authenticated or demo identity
GET /satellites	Cursor, limit, search	Workspace-scoped page and total
POST /satellites	Validated metadata	Created satellite revision
PATCH /satellites/{id}	expected_revision, editable metadata	New revision or 409
POST /imports	Canonical JSON, maximum 2 MiB	Per-record acceptance/rejection report and affected case IDs
GET /cases	Filters, cursor, limit	Stable queue page, total, urgency counts and revision
GET /cases/{id}	None	Current assessment, evidence and case metadata
GET /cases/{id}/reports	Cursor	Source-specific immutable history
POST /cases/{id}/assess	expected_revision, policy_id	Bounded deterministic assessment
POST /cases/{id}/actions	expected_revision, action, rationale, assessment_id	Recorded disposition with updated case
POST /cases/{id}/investigate	assessment_id, optional question	Bounded ADK result or visible template fallback
POST /fleet/compare	Candidate bundle ID and revision	Bounded comparison result or resumable batch state
POST /operations/{id}/continue	continuation_cursor, expected_revision	Process the next bounded batch under a database lease
POST /economics/evaluate	EconomicScenario	Formula result, included/excluded items and source bindings
POST /reentry/evaluate	ReentryScenario	Exposure and conditional damage outputs
POST /communications/simulate	Scenario ID, packet action, expected_revision	Simulation state only
GET /operations/{id}	None	Persisted operation status and available result
GET /activity	after_sequence, limit	Workspace-scoped activity delta
GET /exports/{case_id}	format=json or markdown	Complete versioned evidence packet
POST /workspaces/{id}/reset	expected_revision	Reset only this visitor's sandbox to a seed

All paths except liveness/capabilities require the workspace identity defined by the selected storage mode. A workspace ID in the request is not authorization. Export filenames are server-generated. Reports and scenario inputs are POST bodies, not URLs to fetch.

POST /imports uses ImportBatch: schema_version, expected_workspace_revision and a records array of 1 to 500 canonical input objects, still subject to the total 2 MiB/capacity bounds. Validate the outer envelope first; malformed envelope returns 422. Then validate each record and return an HTTP 200 import summary with

accepted, deduplicated, conflicted and rejected counts plus indexed errors. A wholly rejected but well-formed batch has zero accepted records. Stored accepted records and their summary commit atomically; a storage failure rolls them all back. ImportBatch's full structural schema is useful for strict preflight validation, while the runtime parser retains per-record errors instead of letting automatic whole-body validation discard the entire summary.

Queue behavior

Default page size 50, maximum 100. Sort by urgency ascending (P0 first), then known review deadline ascending with unknown deadline first within its tier, then material concern flag, then case ID. Never sort unknown probability as zero. An analyst can choose alternative display sorts, but the pinned urgent summary remains. Return both filtered result counts and a clearly labeled all-workspace unresolved urgent count.

Pagination cursors carry queue revision and the sort tuple. If time or evidence changes the queue between pages, return 409 QUEUE_CHANGED and ask the client to refresh. This prevents silent skips/duplicates in a moving ranked queue. Snapshot exports can use an immutable queue revision.

Errors and idempotency

Error shape: code, message, request_id, retryable and field_errors. Use 400 for malformed JSON, 401/403 for identity/permissions, 404 for inaccessible or missing objects, 409 for revision conflicts, 413 for oversized input, 422 for schema/domain input errors, 429 for admission limits and 503 for required infrastructure unavailable. An applicable record with scientifically unsupported calculation is a successful domain response with status unsupported, not an invented number or a generic 500.

POST requests that create revisions or expensive operations accept Idempotency-Key. Persist key, request digest and result reference atomically for 24 hours. Same key/digest returns the original result; a different digest returns 409. Authorization is rechecked on replay. Do not use only process memory for this guarantee on Vercel.

Persistence implementation rules

Database storage uses separate immutable input/assessment tables and mutable case pointers. Index workspace_id + urgency + deadline, workspace_id + event_key, report source/revision, dependency digest, and operation status/lease expiry. Use a uniqueness constraint for report identity and idempotency keys. Every table containing user data is workspace-scoped.

Writes use one transaction for revision insertion, current pointer change, invalidation and audit event. A transaction or function crash must not publish an assessment pointing to missing inputs. Store small raw synthetic JSON in the database; do not require an object-storage service for the hackathon. The database/auth provider-specific implementation is finalized in documents 14 and 19.

Cross-field validation and import behavior

Require distinct report object IDs matching the encounter states; interval start < end; unique catalogue IDs; selected object present exactly once; alternatives replacing only that object; candidate IDs unique and different from reserved baseline. State vectors and position covariances use the common epoch/frame declared at the bundle root. The independent covariance model forbids an extra cross-covariance; the supplied-cross-covariance model requires it and a valid full joint matrix. The initial fleet bundle supports independent object errors only; a correlated fleet needs an expanded joint model and is unsupported.

Reentry windows must be ordered; longitude is within [-180,180], latitude within [-85,85], rings are closed and non-self-intersecting, holes are contained, and dateline crossings are unsupported. The combined population/asset point count is at most 1,000, even though each structural array has its own bound. Identical duplicate point IDs deduplicate; conflicting duplicate bodies are rejected. Domain geometry checks supplement the schema and do not guess a corrected polygon.

All monetary terms in a calculation require one currency unless displayed separately. Parse probability JSON numbers through their preserved decimal text for Decimal arithmetic; do not introduce a binary-float rounding step before multiplying money. Store a probability's numeric display and its canonical decimal representation/source reference together. Unknown probability or conditional loss produces an unavailable result rather than zero.

Structural JSON/schema errors reject that record with 422. A structurally accepted report whose complete encounter is scientifically unsupported remains in the ledger with an unsupported calculation finding. A missing encounter is allowed for reported-risk audit, with its missing-evidence findings. Do not discard useful source history merely because the solver cannot run. Report whether a validation error rejected an input or a retained record produced an unsupported result.

Orbital mathematics and deterministic triage

This is a deliberately bounded scientific specification. The implementer must not extend its validity from fixture agreement to arbitrary orbital events. NASA distinguishes screening, assessment and mitigation; ORBIT-TRUST is primarily assessment evidence and review support.¹

Supported encounter model

The input supplies two positions, velocities and position covariances at a common encounter epoch in a common Cartesian inertial frame, plus hard-body radii. initial release accepts GCRF and explicitly synthetic inertial coordinates. Mixed frames, unequal epochs without supplied conversion, unknown covariance basis and undeclared cross-object dependence are unsupported. Do not treat each object's RTN axes as the same frame.

The short encounter assumes straight relative motion, approximately constant projected covariance, independent Gaussian position errors (unless a complete cross-covariance is supplied), and a combined circular hard-body region. Velocity uncertainty is neglected by this particular method and must be declared. Slow, repeating/co-orbital and long nonlinear encounters require another validated method and are not silently forced into this one.

The canonical JSON is the complete-input path. Historical ESA CSV provides reported-risk replay and audit only where field conventions are established. initial release does not claim a complete CCSDS CDM parser or full standard conformance. A future KVN/XML adapter must be a separately tested addition.⁴

Geometry and probability

Let $\mathbf{r} = \mathbf{r}_{\text{secondary}} - \mathbf{r}_{\text{primary}}$ and $\mathbf{v} = \mathbf{v}_{\text{secondary}} - \mathbf{v}_{\text{primary}}$. For a supplied linear segment, the unconstrained time of closest approach relative to the epoch is $\tau = -(\mathbf{r} \cdot \mathbf{v}) / (\mathbf{v} \cdot \mathbf{v})$. Require it to lie inside the declared supported interval; a clipped endpoint is not a completed encounter for this 2D calculation. At a supplied TCA, check orthogonality with $|\mathbf{r} \cdot \mathbf{v}| \leq 1e-8 \max(1, |\mathbf{r}||\mathbf{v}|)$; otherwise recompute within the supported synthetic segment or mark the direct-TCA input inconsistent.

For independent errors, $\mathbf{C}_{\text{rel}} = \mathbf{C}_{\text{primary}} + \mathbf{C}_{\text{secondary}}$ in the common frame. With supplied cross covariance $\mathbf{C}_{\text{ps}}$, use $\mathbf{C}_{\text{primary}} + \mathbf{C}_{\text{secondary}} - \mathbf{C}_{\text{ps}} - \text{transpose}(\mathbf{C}_{\text{ps}})$. The complete joint covariance must pass validity checks; a random cross term must not be accepted because the projected result happens to look positive.

Construct the encounter-plane basis deterministically. Set $\mathbf{n} = \mathbf{v}/|\mathbf{v}|$. Select the Cartesian unit axis least aligned with $\mathbf{n}$ (ties x, then y, then z); set $\mathbf{e1} = \text{normalize}(\text{axis cross } \mathbf{n})$, $\mathbf{e2} = \mathbf{n} \text{ cross } \mathbf{e1}$. Project the closest-approach vector and relative covariance onto rows $\mathbf{e1}$, $\mathbf{e2}$ to obtain mean $\mathbf{m}$ and 2 by 2 covariance $\mathbf{S}$. Combined radius R is the sum of the two supplied radii, never a diameter or radar cross section substituted without evidence.

P_c is the integral of the 2D normal density with mean $\mathbf{m}$ and covariance $\mathbf{S}$ over $\mathbf{x}^* \mathbf{x} + \mathbf{y}^* \mathbf{y} \leq R^2$. Implement deterministic Gaussian quadrature over polar coordinates, including the radial Jacobian. Compare a 64 radial by 128 angular grid with 128 by 256. Accept numerical convergence only if the absolute difference is $\leq \max(1e-12, 1e-7 \text{ times the refined } P_c)$. Use the refined value. If it fails, try 256 by 512 once; otherwise return numerical_nonconvergence. Precompute quadrature nodes and weights from a tested algorithm/library, not hand-entered approximate constants.

Use a Cholesky solve and log normalization to avoid explicit matrix inversion when evaluating the density. If floating-point underflow yields zero without an error bound, report below_computable_precision with null authoritative P_c ; do not claim impossible collision. Positive tiny values can be displayed using scientific notation. Never round a probability before evaluating policy.

These tolerances are numerical engineering defaults, not physical uncertainty estimates. Relative speed below 1 m/s, nonpositive projected covariance, projected condition number above $1e6$, invalid radius, unknown units or missing applicability declarations are unsupported in initial release. Raising these limits requires new reference cases and an explicit method version.

Covariance checks

Require symmetric input within $1e-10 \max(1, \text{infinity norm of } C)$ in the declared units. Material asymmetry is invalid. Do not silently repair covariance. The initial method accepts positive-definite matrices; singular positive-semidefinite matrices may be physically meaningful but are unsupported by this solver. Check finite entries, nonnegative variances and correlations in $[-1, 1]$. A correlation of 1 can lead to singularity and must not be passed to inversion.

When standard deviations and correlations are explicitly defined, reconstruct $C_{ij} = \rho_{ij} \sigma_i \sigma_j$. Covariance scale a multiplies variance, so standard deviation scales by $\sqrt{a}$. Store the original input and the declared scenario transform. Algebraic validity does not demonstrate covariance realism; that requires empirical orbit-error or residual evidence.⁵

Checked reference fixture

For $m=(100,0)$ m, $R=10$ m and $S=\text{sigma squared times identity}$, the following values were independently checked by Cartesian and polar integration during preparation of the preceding playbook. They are reference facts for a synthetic encounter, not an operational engine benchmark.

Combined sigma in m	Pc
20	8.71274018586874e-7
50	0.002733592576274527
200	0.0011025180765032241
1000	4.9749386433385265e-5

Use an independent SciPy noncentral-chi-square reference for isotropic cases and a separate Cartesian integration for general cases. Test the exact zero-miss isotropic result $P_c = 1 - \exp(-R^2 / (2 \sigma^2))$. Full test tolerances are in document 12. Probability dilution is an established phenomenon, not a novelty claim.⁶

Triage policy version demo-2.0

Review threshold $\theta = 1e-4$, immediate-review slack = 3600 seconds. These are illustrative mission policy choices, not universal operational instructions. Material concern is true when an applicable provider or supported recomputation crosses θ , a declared applicable scenario crosses θ , or a mission-specific supplied override requires review. Missing information needed to classify concern also requires review, but does not assert high collision probability.

Tier	Condition	Display
P0	Required unresolved review and current time is later than its known review deadline	Overdue review; persistent escalation
P1	Required review and deadline unknown, or slack ≤ 3600 seconds	Review now; exact reason and remaining time
P2	Required review and known slack > 3600 seconds	Scheduled review; owner and checkpoint
P3	No material concern, usable applicable evidence and stable outcome across the declared supported scenarios	Monitor; next checkpoint and reopening conditions

An optional cosmetic field is not material. Unusable covariance needed for a recomputation cannot become $P_c=0$. Valid reported risk may still support a review requirement while recomputation is unsupported. A provider-neutral label does not imply independent observations.

The review deadline is a supplied latest feasible command opportunity minus the supplied review/coordination allowance. Do not invent it from TCA alone. Fixture: command opportunity 14:30 UTC, review allowance 90

minutes, review deadline 13:00 UTC; at 12:20 UTC, slack is 40 minutes and a required case is P1. At exactly the deadline it is P1; after it, P0.

TCA passing does not prove no collision. The separate event_phase becomes post_encounter_pending_verification, and unresolved required review is retained until a human records an appropriate disposition. This is not an extra workflow stage or an automatic case closure. Historical post-TCA rows are excluded from predictive evaluation, separately from this workflow rule.

Evidence audit rules

The canonical evidence_metadata distinguishes orbit-solution epoch and last-observation time from the propagated encounter epoch. Age at the review clock is now minus the corresponding supplied solution/observation time. Never calculate orbit age from the future TCA state epoch, or infer it from ESA observation offsets with unresolved sign conventions.

Demo defaults require solution and last-observation metadata no older than 21,600 seconds; values over that limit produce material stale-evidence findings. This six-hour rule is a fictional configurable review policy, not a universal orbit-quality threshold. Missing required metadata, unknown maneuver information, or an unverified covariance-realism basis produce material missing-evidence findings. Future observation/solution timestamps more than five seconds beyond the scenario clock produce an inconsistency; the five-second tolerance is a declared demo input allowance. The synthetic_assumption covariance basis is usable only in synthetic mode, with a persistent notice that realism has not been established for real objects. Operator_declared means a supplied assertion requiring its source reference, not independent calibration.

Use stable finding codes: MISSING_FRAME_EPOCH, INVALID_COVARIANCE, UNSUPPORTED_ENCOUNTER, STALE_SOLUTION, STALE_OBSERVATION, MISSING_OBSERVATION_METADATA, FUTURE_EVIDENCE_TIME, COVARIANCE_REALISM_UNVERIFIED, MANEUVER_CONTEXT_UNKNOWN, SOURCE_REVISION_CONFLICT, RISK_THRESHOLD_EXCEEDED, SCENARIO_CLASSIFICATION_CONFLICT and DEADLINE_UNKNOWN. Each finding includes object/report IDs, source JSON pointer, observed value or missing marker, rule version, material flag and host-rendered text. A deadline gap affects timing classification, not the numerical integration.

Optional supported covariance sensitivity scenarios are declared inputs, not an exhaustive search or a new independent observation. For the four-sigma reference, explicitly identify the selected alternative(s) and compare their policy classifications. Changing sigma from 200 to 1000 lowers Pc below the demo threshold; that alone cannot clear an urgent floor or resolve a stated uncertainty conflict.

Escalation and acknowledgement

Store proposed_urgency separately from unresolved_alert_floor. Automatic updates can raise urgency immediately. A lower recommendation does not silently clear an unacknowledged P0/P1 alert. An analyst action must reference the latest assessment and provide a rationale to clear that floor. Closing a case requires a disposition and latest-version acknowledgement. New material evidence reopens it as Reviewing; duplicates and cosmetic metadata do not.

Capacity limits never hide required reviews. Top-K evaluation is an analysis metric, not permission to delete overflow. Show overload counts and an escalation instruction when required reviews exceed configured capacity. Economic value cannot lower a required safety review.

Source notes

1. NASA, [Spacecraft Conjunction Assessment and Collision Avoidance Best Practices Handbook](#), 2023; NASA CARA, [Close Approach Risk Mitigation](#).
4. CCSDS, [Conjunction Data Message, 508.0-B-1 with corrections](#).
5. NASA NTRS, [Covariance realism research record](#).
6. Kayhan Space, [Risk metrics](#).

Fleet comparison and communication simulation

Candidate input and coverage

Accept an immutable candidate set containing a baseline and at most three alternatives for one selected spacecraft, with up to 20 catalogue objects in the public demo. The bundle supplies common-epoch Cartesian state and position covariance for each object, the supported short encounter interval, and candidate planning metadata. All objects are synthetic in the public demo. Alternatives replace only the selected object's state within that interval.

The first demonstration uses three objects with locally linear short arcs. It does not simulate gravity over a complete orbit or derive the candidate from a burn. Candidate positions encode hypothetical supplied outcomes. Label the horizon prominently: "All 3 supplied objects assessed over this 10-second synthetic encounter segment." Do not shorten this to "Fleet safe."

Evaluate every unordered pair for each candidate in the small catalogue. A 20-object catalogue has 190 pairs. Reuse unchanged pairs, but include their actual results in coverage. Use the sorted pair key, source revisions, scenario/method versions and horizon in the cache key. Do not add an unvalidated pruning step merely to claim optimization.

Comparison logic

For each pair, compute supported geometry/ P_c and apply the same supplied policy. Return a matrix of candidate by pair containing distance, P_c if supported, material findings and evidence links. A candidate is `blocked_by_demo_policy` if any applicable pair violates the threshold or a supplied hard feasibility constraint fails. If a required pair/input is unsupported or coverage is incomplete, the result is `review_required_incomplete`, never a pass. Otherwise it is `passes_demo_checks`.

The selected object's maneuver capability, command opportunity, planning allowance and candidate feasibility flag must be supplied. A flag proves only that the input asserts feasibility. Expose this distinction as `feasibility_basis=supplied_assumption`. In initial release there is no onboard propulsion, thermal, attitude or fuel-system validation. Analyst approval records a review choice; it does not promote supplied assumptions into independently verified physics.

Among candidates that pass the declared checks, sort by supplied `delta_v_m_s` ascending, then supplied response cost, then candidate ID. Unknown cost or delta-v sorts last and remains missing. Do not call the selected candidate globally optimal: it is preferred among the finite supplied alternatives under this comparison rule. Compare safety feasibility before money.

Synthetic demonstration

The fixture supplies A, B and C with plausible instantaneous orbital-scale positions/velocities, but explicitly models only a short linear segment. Baseline A-B has small miss distance and material concern. Candidate one changes A's supplied position to reduce the A-B concern but coincides with C at closest approach. Candidate two moves A away from both within the supplied segment. The fixture includes independent expected pair probabilities and decision labels.

This demonstrates fleet conflict checking, not automated maneuver synthesis. The UI should let judges switch supplied alternatives, see the affected pair and inspect assumptions. Arbitrary dragging on an orbit graphic must not generate a purportedly executable maneuver.

Optimization boundaries

Use Rayon inside the Rust core for independent pair jobs, with a bounded worker count. Separate algorithm improvements from language/runtime improvements in the benchmark. Compare the same inputs, tolerance, work and cache state. Report cold-start, uncached compute and cached response separately. Ordered output and stable tie-breaking must not depend on thread completion order.

Do not combine pairwise probabilities into a single “probability of any fleet collision” using a sum or independence formula without a joint event model. Repeated warnings can describe the same encounter, and a single spacecraft loss can appear in several pairs. The initial release interface displays pairwise results and unique case counts. Full-catalog screening, collision cascades and Kessler syndrome simulation are future separate capabilities.

Communication state machine

Communication runs entirely in a sandbox with scripted packets. States are draft, authorized_for_simulation, waiting_for_contact, sent, acknowledged, execution_reported, expired, failed and cancelled. Only supplied simulation events can move sent to acknowledged and then to execution_reported. The interface uses those exact distinctions. There is no real satellite, radio network or external message connector.

A packet records case/assessment/candidate IDs and revisions, creation/expiry time, channel label and idempotency ID. A new material assessment invalidates an unsent old packet. A stale packet cannot be authorized. The simulator rejects sends outside the supplied contact window and marks packets expired at their explicit expiration. Duplicates return the prior packet state; acknowledgement of one packet cannot acknowledge another.

Feasibility uses supplied end-to-end delays: wait until contact + uplink time + acknowledgement allowance + required action lead time must fit before the supplied action deadline. None of these is derived from Groq token speed or HTTP response time. When input is missing, display timing_unknown. The application cannot promise to protect a spacecraft simply because its calculations finish quickly.

Require the complete uplink interval to fit inside the contact window. If acknowledgement uses the same supplied link, its allowance must also fit that window. Packet expiration is inclusive: at now $\geq$ expires_at it is expired. A sent packet with no acknowledgement becomes failed at its acknowledgement timeout and cannot be labeled executed without a valid subsequent simulation event sequence. Cancelling an unsent draft/authorized/waiting packet sets cancelled. A terminal packet is immutable; reset creates a fresh packet ID. The fixture communications.json exercises these distinctions using its own stated clock, separate from the default queue clock.

Failure cases

Show delayed or absent acknowledgement, lost packet, expired candidate, unavailable contact, impossible supplied planning deadline and conflicting concurrent candidate selections. An analyst can cancel a draft or reset the simulator; this changes only the current workspace. Record all transitions in the audit ledger. In real future integration, the protocol, authorization, command signing and flight-control validation need their own independently reviewed specification.

Reentry exposure and financial scenarios

Reentry input boundary

The reentry module receives a supplied footprint, never a conjunction report alone. Inputs are a GeoJSON Polygon or MultiPolygon in WGS84 longitude/latitude order, a declared scenario time window, source/revision, and exposure layers. initial release supports regional polygons that do not cross the antimeridian, latitude within -85 to 85, and valid non-self-intersecting rings. Reject unsupported global/dateline geometry visibly instead of drawing the long way around Earth.

The canonical exposure layers are arrays of synthetic point records: population sample points carry an integer `represented_population`; asset points carry `asset_id`, `asset_type` and `replacement_value` as an amount/currency object. The footprint is a GeoJSON geometry object; a full FeatureCollection adapter is optional future work. The footprint test includes its outer boundary and excludes the interior of holes. A point on a hole boundary counts as exposed under a documented conservative convention. Deduplicate identical asset IDs and population sample IDs before summation; reject conflicting duplicate records. Overlapping MultiPolygon parts do not double count.

The initial implementation uses a tested point-in-polygon library and explicit boundary tests. It does not convert geographic degrees into metres using one constant. No area-based population estimate is claimed: these are point representations, and the UI states “population represented by included sample points,” not a validated census-level casualty estimate.

Outputs

Return included/excluded point IDs, exposed represented population, exposed asset count, exposed replacement value grouped by currency, unavailable-value count and source-layer coverage notes. Missing asset value remains unknown, not zero. Do not total INR and USD without an explicit scenario exchange rate; default show separate currency totals.

Terrain/land-cover labels provide context. Flat ground is not automatically lower consequence; settlement, critical infrastructure, fragment survival and impact energy can dominate. In the synthetic sandbox, one footprint covers more represented people/assets than another. They are alternative supplied footprints, not locations the system has proved reachable. A low-exposure footprint is not a recommended real disposal target.^{2, 3}

Optional conditional damage

Only enable monetary damage when the scenario supplies, for each included asset, a `probability_of_damage` conditional on this reentry scenario and a `mean_loss_fraction` conditional on damage. Then expected damage for that asset = `probability_of_damage` times `mean_loss_fraction` times `replacement_value`. The same impact can damage multiple assets, so expectation of the sum is the sum of expectations without requiring independent damage events. This does not establish the probability that any damage occurs.

Keep excluded/missing vulnerability items visible. Sum only known terms and label a partial estimate if any included asset lacks assumptions. Do not infer damage probability by dividing polygon area or using TerraSight's safety score. Do not monetize people, infer casualties from point counts or produce a legal liability estimate.

Orbital economic model

The owner supplies disjoint conditional loss components: spacecraft replacement, launch/replacement service, unrecovered mission revenue, service interruption and other documented costs. Avoid including the same lost revenue in both mission revenue and interruption. Separate insured recovery if modeled; default insurance, legal liability, environmental monetization and third-party loss are excluded.

Conditional loss L is the sum of included monetary terms in one currency. A single-event expected loss is P_c times L only when P_c applies to the same event and L describes loss conditional on that event. If P_c is unsupported, conflicting beyond a declared scenario or numerically unresolved, report the financial result as unavailable or provide separate labeled scenarios, not a blended estimate.

For a fixed comparable conditional loss, estimated net benefit of response = $(P_{c_before} - P_{c_after}) \times L - \text{response_cost}$. With different conditional-loss assumptions, use $P_{c_before} \times L_{_before} - P_{c_after} \times L_{_after} - \text{response_cost}$. Response cost can include supplied operations labor, fuel/lifetime opportunity cost and interruption caused by the response, without counting the same cost twice.

Use decimal arithmetic and round only the final display to the currency's minor unit. Stored values retain sufficient precision for reproduction. Input monetary values are decimal strings. Reject negative probabilities, negative declared costs, mixed currency arithmetic and nonfinite numeric fields. A negative net benefit is displayed, not clipped to zero. It does not cancel a required review or hard safety policy.

Reference economic case

Entirely synthetic inputs: conditional loss INR 1,000,000,000; P_{c_before} 0.0001; P_{c_after} 0.000001; response_cost INR 20,000. Expected loss before is INR 100,000; after is INR 1,000; gross expected-loss reduction is INR 99,000; net benefit is INR 79,000. Displaying INR 1,000,000,000 as “saved” would be incorrect.

One warning's updates do not create separate economic events. Do not sum a spacecraft's full loss repeatedly across overlapping encounters. The initial release financial view is per scenario/case; portfolio aggregate benefits are out of scope until event dependencies and mutually exclusive loss pathways are modeled.

Uncertainty presentation

Show low/base/high user-supplied loss assumptions as scenarios, not confidence intervals. If probabilities come from different applicable covariance alternatives, show each probability-economic pairing with its provenance. Do not label the maximum among a handful of scenarios as the mathematical maximum possible risk.

Actual money saved is a later empirical/causal claim requiring a defensible counterfactual and observed costs. Hackathon metrics use “modeled expected-loss change” and “estimated review effort saved.” The reentry and orbital economic branches must never multiply orbital P_c directly by all asset value under a ground footprint: the intermediate breakup/reentry/impact probabilities are absent.

Source notes

2. ESA, [Reentry background](#).

3. ESA, [Re-entry safety](#).

ADK orchestration and output validation

Purpose and model choice

The agent's useful task is bounded investigation: decide which available evidence history or declared scenario needs inspection and organize a review packet. Numeric integration, deadline arithmetic, financial math, authorization and policy remain deterministic. Do not generate decorative agent traces from a scripted animation.

Use Google ADK Python with its LiteLLM connector and Groq model ID `groq/openai/gpt-oss-20b`. The raw Groq model ID is `openai/gpt-oss-20b`; the first `groq/` prefix belongs to LiteLLM routing. The server reads `GROQ_API_KEY` from its secret environment. Verify the model in the owner's account and run a real tool-call/schema spike before locking the provider adapter. The earlier Finora Llama default is not assumed available to a new free account.^{7, 8, 9, 10}

Groq is selected for remote inference without the team's own GPU and documented function/schema capabilities; local latency and account quotas still need measurement. Current Groq documentation lists tool use and strict structured output for the selected model, with those features used in separate calls. The selected model does not require parallel model tool calls: deterministic batch execution provides the numerical parallelism. No built-in web search, code execution, remote MCP or provider-managed actions are enabled.^{8, 9, 10}

The supplied ADK lecture is the component reference, not a mandate to use every component. Use `LlmAgent`, `FunctionTool`/custom deterministic tools, `SequentialAgent` or a small custom `BaseAgent` workflow, typed output, callbacks and versioned state. No A2A, general long-term memory or multimodal model is needed. Its warning about mixing tools and schema is addressed by separate stages. The supplied deck and a text extraction are included under `reference/`.

Roles and orchestration

The overall `CaseReviewWorkflow` is deterministic. It verifies the assessment snapshot and executes the required calculations before invoking any LLM. Domain roles `EvidenceAudit`, `FleetComparison`, `ConsequenceAnalysis` and `ReentryExposure` are explicit deterministic services/tools. Only `Investigator` and `PacketFormatter` invoke the model. This is a multi-role system, not six independent language models.

`Investigator` receives a compact record of the current assessment, source identifiers and available scenario IDs. It may call at most three allowlisted tools in total: `inspect_event_history`, `inspect_evidence_finding`, `evaluate_allowed_scenario` or `read_candidate_comparison`. Tool arguments are only existing IDs, never a new orbit, probability, cost or arbitrary URL. Bound each result to 4 KiB of selected fields; full artifacts remain in the ledger.

Use at most two tool-response rounds followed by one formatter request. Maximum normal provider requests: three, plus one retry shared across the entire run, never a retry per agent/tool. If no further tool is appropriate, proceed directly to formatting. Budget at most 1,200 input tokens per provider request and 400 output tokens per call, with an overall 6,000-token reservation including retry. Log actual token use and release unused reservation. These are application defaults to adjust downward for the actual account, not Groq entitlements.

Apply a 30-second total agent wall-time budget including tools and any retry. Each provider request has a timeout of at most 8 seconds or the remaining budget, whichever is smaller. Before dispatch, reserve enough remaining time for host validation and persistence. A fourth call is permitted only if both token and time budgets still allow it. Long scenario computation should already exist as a saved bounded operation; an agent tool must not bypass the numerical chunk limits.

The formatter has no tools. It returns a typed selection: `assessment_id`, `ordered_finding_ids`, `selected_fact_ids`, `summary_template_code`, `request_template_code` or null, and `question_to_analyst_code` or null. All fields are required; nullable values are explicit; `additionalProperties=false`. Use ADK `output_schema` if the verified connector forwards the needed schema reliably. If that spike fails, call Groq's strict-schema formatter within an ADK custom stage and validate with the same Pydantic model. Do not silently remove schema validation or report a mocked provider response as live.

Authoritative text assembly

The initial release accepted packet is rendered by the host from validated template codes and fact records. The model chooses evidence order and permitted templates; it does not write arbitrary factual sentences into the authoritative packet. This closes the gap left by merely checking digits in free text. Numeric facts, next action and urgency are always host-owned.

Allowed summary templates include REVIEW_DUE_TO_RISK, REVIEW_DUE_TO_MISSING_EVIDENCE, REVIEW_DUE_TO_CONFLICT, MONITOR_SUPPORTED and CALCULATION_UNSUPPORTED. Allowed request templates include REQUEST_STATE_COVARIANCE, REQUEST_FRAME_EPOCH, REQUEST_MANEUVER_CONTEXT and REQUEST_DEADLINE. Their availability is computed from findings. A template cannot be selected if its prerequisite finding is absent.

An example accepted packet says: “Review is required because the applicable reports support different classifications. The current review deadline is {deadline_fact}. Request an updated {missing_field_names} with source and epoch before {request_deadline_fact}.” Placeholders come from host records. If a fact is unavailable, select the corresponding missing-data template rather than inventing a value.

An optional free-text “AI wording draft” is outside initial release. Do not add it as a shortcut around these requirements. The product still demonstrates agentic choice through different allowed investigation sequences and evidence selection; evaluate whether that choice is useful against fixed routing.

Guardrails and failure handling

Before any tool call, verify run/workspace identity, expected input digest, allowed tool name, schema, argument IDs and remaining budget. Report comments, names and imported notes are untrusted text. They cannot change tool policy, trigger network calls or become system instructions. Mask contact information from model inputs by default; public data is synthetic.

After each tool result, verify provenance and status. Before publishing, ensure every cited finding/fact exists in the same snapshot, selected templates match facts, and no required finding is omitted. The deterministic required-finding list is included even if the model fails to select it. A schema-valid but wrong selection is rejected. New source data marks the run stale, and its packet is not displayed as current.

On malformed response, forbidden tool, invalid reference, provider 429, timeout or missing key, return the deterministic packet with generation_mode=template_fallback and a plain explanation. Retry only if it fits the single retry and remaining time/token budget. Respect Retry-After; if waiting would exceed budget, stop. Never switch to a paid provider or claim the fallback is Groq output.

Quotas and observability

Account limits are shared across organization requests and tokens; inspect actual values rather than copying a fixed public table. Maintain atomic database reservations for the chosen model and account. Default global active model runs=1, one concurrent run per user, three investigate runs/minute/user, ten runs/day/anonymous user and fifty runs/day/persistent user. Provider request counts are separately bounded within each run. Global spend remains disabled. Administrator limits cannot exceed the verified provider allowance. Seed-page views and ordinary calculations make no model calls.⁹

Record model ID, provider, prompt version, accepted tool names/arguments, source record IDs, stage durations, tokens, retry cause and validation outcome. Do not store hidden chain-of-thought, API keys or full provider headers. Show “what ran and what it returned,” not claims about the model's internal thoughts.

Dependency note: official ADK documentation identifies LiteLLM 1.82.7 and 1.82.8 as compromised. Exclude these releases, choose a currently patched compatible version and lock it in M0. This is a dependency-selection requirement for the build, not evidence that this workspace has installed those versions.⁷

Source notes: 7. Google ADK, [LiteLLM integration](#). 8. Groq, [Supported models](#). 9. Groq, [Rate limits](#). 10. Groq, [Tool use and Structured outputs](#).

Mission console interaction specification

Visual system

Use the owner-selected clean mission console: dark navigation, light work panels and restrained status colors. Sidebar background #111827, work canvas #F4F6F8, panels white, primary text #111827, secondary text #475569, borders #D5DCE5, primary action #1D4ED8. P0/P1 use red text/icons with pale backgrounds, P2 amber and P3 blue/neutral. Green is reserved for a completed software check, never general spacecraft safety. Verify actual contrast in rendered controls.

Use a local/system sans-serif font, 16 px body, 14 px dense table text, 24-28 px page title and 12 px metadata only where legible. Default line height 1.45. Use 8 px spacing increments, 8 px panel radii and minimal shadow. No animated starfield, background music or full-screen 3D globe. The central visual is an evidence timeline or pair comparison that explains a decision.

At desktop widths use a 224 px sidebar and a content width that fills the remainder with 24 px padding. Below 900 px use a collapsible navigation drawer. At 390 px the primary action, case reason and urgency remain readable; wide scientific tables scroll inside their region, not the entire document. Reduced-motion preference disables nonessential animation.

Navigation and screens

Home offers a short project statement, Start demo, Sign in and a clear synthetic-data notice. A signed-in user returns to their saved workspace. Persistent identity uses GitHub OAuth; Try demo uses Supabase anonymous sign-in with a private workspace. The static shell fetches session data only on the client.

Queue is the default work screen. Top row shows P0/P1 count, evidence issues, pending review and last update. The table columns are urgency, case/object pair, review deadline/slack, reported Pc, evidence state, main reason, assigned analyst and stage. Unknown values use an em dash with “not supplied” accessible text. Full queue counts remain visible beyond pagination. A persistent strip highlights any unacknowledged urgent cases and capacity overflow.

Case detail shows “Why this needs attention” first, followed by deadline, current evidence and recommended review step. Tabs: Evidence, Reports, Response comparison, Economics, Activity. The source timeline distinguishes creation time from reception time and provider streams. A side panel holds assignment, notes, acknowledgement and review disposition. Disable stale-version actions until refreshed; explain the conflict.

Evidence view shows passed/failed/unsupported checks with expandable source fields. An encounter-plane plot shows mean and uncertainty geometry only when inputs support it. Axes have units, source revision and “schematic” label where applicable. Do not render an Earth map from missing absolute coordinates.

Fleet view lists exactly which objects and time interval were assessed. Rows are pairs; columns are baseline and supplied alternatives. Selecting a cell opens its evidence. A candidate that helps A-B while creating an A-C concern displays both changes side by side. Candidate recommendation text says “Preferred among supplied options under demo policy.”

Satellite registry supports name, operator label, maneuver capability, contact-window metadata and economic assumptions. Orbital source records are read-only imports; correcting a scientific input creates a new revision with reason. Do not let a simple edit form silently rewrite a provider report. Archive a satellite only when retaining references; historical case records are never cascade-deleted by a metadata action.

Reentry Sandbox is a separate navigation item with a permanent banner: “Supplied footprint scenario. No reentry trajectory prediction.” Show footprint alternatives, exposure points, legend, coverage and included/excluded counts. Base maps are optional local outlines; no paid map key or runtime map-tile dependency is required. Map selection never sends an orbital maneuver.

Activity shows recorded tool execution, calculation completion, report import, assignment and review changes. It omits raw hidden reasoning. Settings contains display currency, account identity, data export/import and demo reset. Policy edits require a new version with visible history; a public anonymous visitor can change only their isolated sandbox.

States and accessibility

Every screen has loading, empty, error, stale and partial-data states. A slow model does not replace already available numerical content with a spinner. Provider failure has a small explanation beside the template packet. Infrastructure failure prevents writes and offers retry; it cannot show a false successful save.

Use semantic tables with headers, labeled inputs, visible focus, keyboard-accessible dialogs and action buttons, and an accessible text alternative to every plot/map. Do not communicate urgency by color alone. Avoid automatic focus jumps when queue order changes; announce update counts and let the user refresh. Auto-advancing demo time is off by default.

Destructive sandbox reset requires one in-app confirmation naming that workspace. Editing costs, adding satellites and notes are ordinary reversible actions. A review disposition asks for a rationale because traceability is the product requirement, not because the AI is requesting permission to answer.

Reuse constraints

Adapt Sajawat's admin component patterns and permission-aware detail interactions, replacing commerce terminology and its 100-record client board assumption. Adapt TerraSight's spatial presentation, not its safety score. Finora's execution trace can inspire the activity view, but do not reuse unvalidated confidence badges or theatrical timing as measured latency.

Data sources and reproducible fixture catalogue

Three distinct data modes

Public demonstration mode accepts synthetic canonical JSON and seeded examples only. Local research mode can import the ESA historical training archive for reported-risk replay. A future operator mode needs separate data agreements and complete validated source adapters. A provenance label is visible at workspace, case and export levels.

The ESA archive was downloaded and inspected during prior research. It contains 162,634 rows, 103 columns, 13,154 events and 19 mission IDs; SHA-256 is

68362fe5629cc80f17291f2d73f733bf4e922675e37b91a8ee79afadb46f3edc. Download link: https://kelvins.esa.int/media/public/competitions/collision-avoidance-challenge/train_data.zip. The derived audit is included under reference/; the 87.7 MB archive is not bundled. ¹⁵

Historical adapter

Preserve all columns in a raw research record. Normalize documented risk as log10 Pc; retain the raw value. The 67,240 rows at -30 have unresolved floor/sentinel semantics and cannot establish meaningful precision at exactly 1e-30. All max_risk_estimate values are negative; exclude that field from risk/policy math until its encoding is verified. Observation-offset sign conventions are unresolved; do not label them absolute orbit age.

Exclude 391 negative time_to_tca rows before prospective replay selection. Then select events with at least one input at time_to_tca >= 2 days and a latest pre-TCA record with 0 <= time_to_tca < 1 day: the inspected archive gives 8,287 eligible events. An earlier audit key reported 7,962 because it excluded whole events containing any negative row; that key is superseded and is not the intended cohort.

Sort each event by decreasing time_to_tca for replay. Split by whole event, never by individual report, and do not use the last recorded risk as an actual collision label. With no absolute dates, do not claim a calendar-time train/test split. Freeze deterministic event-based calibration/holdout partitions and report mission composition. This pack does not require ML training.

Some position and velocity covariance/observation fields are missing. Suitability is calculation-specific, but the archive lacks the full geometry/frame/mission information required for our canonical complete-encounter path. Its last-observation metadata, max-risk fields and velocity gaps must not be silently guessed into a validated orbital assessment.

The data page and challenge rules do not establish unrestricted redistribution or commercial reuse for this project. Do not publish the raw archive, private operator data or a derived extract in the public demo without resolving applicable terms. This pack includes only source links, aggregate audit facts and original fictional fixtures. ¹⁵

Bundled fixtures

File	Purpose
encounter_reference.json	Complete synthetic encounter with the four checked sigma alternatives
fleet_candidates.json	Three-object short-segment comparison with supplied A alternatives
triage_cases.json	Clock/deadline/materiality cases and expected tiers
economics.json	Exact decimal arithmetic reference and missing/invalid scenarios
reentry.json	Regional polygons and fictional population/asset points
communications.json	Contact timing, acknowledgement and packet-state simulation

File	Purpose
report_update_cases.json	Duplicate, out-of-order, source conflict and stale-assessment expectations
adversarial_cases.json	Forbidden tool calls, prompt injection and citation/selection failures

Expected values are test oracles for the specified model and fixtures, not stored production answers to display instead of computing. The next implementer must implement the calculation and compare against them. The packet-generating host can use templates, but fake numerical outputs or hardcoded candidate outcomes do not satisfy the numerical milestone.

Fixtures containing malformed values are negative tests, not examples of accepted input. Their expected_rejection field distinguishes schema errors from domain unsupported outcomes. No seed contains a real person's contact data, actual spacecraft command or secretly copied financial record.

Default demo seed recipe

Create three synthetic satellite metadata records A/B/C and an audit case from the sigma-200 Report input. Give that report source_revision=r1 and a matching reported_pc from its reference value. Insert a later r2 at 12:10 UTC with sigma-1000 covariance, its matching lower reported Pc, and secondary_last_observation_at=null. Preserve r1 in history. At the fixed 12:20 clock, the current case requires review for missing evidence and retains the prior unacknowledged urgent floor. The model can inspect history and request observation/state context. Do not automatically call the old covariance a currently applicable scenario unless the scenario set explicitly declares it.

Load fleet_candidates.json as a separate candidate-set scenario and show all nine candidate/pair cells. Load economics.json as the explicitly separate financial reference and reentry.json as the separate ground-exposure reference. The sigma alternatives are independent calculation fixtures; do not bulk import all four unchanged source revisions as if they were sequential reports. The report-update fixture describes deliberate duplicate/conflict cases separately.

For the hero recording, retain one additional deliberately conflicting source-revision example if demonstrating SOURCE_REVISION_CONFLICT. Missing data, an older warning and an actual same-revision conflict are different findings; the narration must match the displayed case. Seed reset restores the same fictional clock and source history using fresh runtime IDs.

Generation and scaling

For >100-case pagination tests, generate 205 explicit unique synthetic case IDs from the reference family with recorded seed=20260911, varied supplied deadlines and 7 known P0/P1 cases beyond insertion index 100. These are UI/load fixtures and not 205 independent scientific encounters. For numerical load testing, disclose the number of repeated versus distinct geometries and disable cache in the uncached pass.

Public maximums: 20 satellites, 500 cases, 2,000 report revisions, 3 response alternatives plus baseline, 5 reentry scenarios and 1,000 exposure points per workspace. Reject larger public bundles with a limit message and recommend local research mode. These limits bound costs and validation effort; they are not platform capacity claims.

Keep archive downloads explicit and outside server request paths. Imported filenames are metadata, not filesystem destinations. Public JSON uploads are at most 2 MiB, contain no embedded archives and do not trigger arbitrary URL downloads. Export includes schema/provenance, input hashes, assumptions and method version so another installation can reproduce it.

Source notes

15. ESA Kelvins, [Collision Avoidance Challenge data and Rules](#).

Test strategy and acceptance catalogue

These are tests to implement, not tests already passed by ORBIT-TRUST. The package itself has separate document/fixture QA. Record every application test result with commit, environment, dependency versions and actual output in BUILD_STATUS and the release report.

Numerical and policy checks

ID	Scenario	Acceptance
T01	Four isotropic reference encounters	abs error $\leq \max(1e-12, 1e-7 \text{ times reference } P_c)$
T02	Zero miss, isotropic covariance	Match analytic $1 - \exp(-R^2/(2 \sigma^2))$ within T01 tolerance
T03	General positive-definite anisotropic fixtures	Match independent Cartesian reference; report convergence
T04	Rotate all states/covariances by one orthogonal matrix	P_c invariant within T01 tolerance
T05	Swap primary/secondary consistently	Same distance/ P_c and equivalent policy
T06	Common position/velocity translation	Relative encounter result unchanged
T07	Invalid/singular/asymmetric covariance, missing frame, nonfinite fields	Reject or unsupported as specified; no zero-risk substitute
T08	Very low relative speed, endpoint-only closest approach, unsupported nonlinear declaration	Unsupported method state
T09	Nonconvergent/underflow numerical cases	Explicit failure or precision status, no definitive zero
T10	Increase hard-body radius with same valid inputs	P_c nondecreasing within tolerance
T11	Sigma 200 versus 1000 reference alternatives	Lower P_c under greater uncertainty does not erase conflict
T12	Deadline +40 minutes, >60 minutes, unknown, exact and passed	P1/P2/P1/P1/P0 as specified
T13	Missing optional cosmetic field versus missing material evidence	Only the material gap affects required review
T14	Downgrade after urgent alert	Alert floor remains until latest-version acknowledgement
T15	New material evidence on a closed case	Reopens and invalidates stale packet

Use a separate implementation for numerical expected values. The reference Python code may use SciPy `ncx2` for isotropic cases and adaptive Cartesian quadrature for general cases. Sharing input parsing is acceptable; using the Rust function to generate its own sole oracle is not. Record deterministic sum ordering and tolerance across thread counts; do not require unsupported cross-CPU bitwise floating-point identity.

Workflow, security and integration

ID	Scenario	Acceptance
T16	Duplicate and reordered reports	No false new scientific revision; source history preserved
T17	Same source revision with different body	Conflict retained, no overwrite
T18	Same event key with different object pair	Rejected association

ID	Scenario	Acceptance
T19	205 cases with urgent records after index 100	All urgent cases accessible; counts accurate
T20	Queue changes during pagination	409 refresh; no silent skip
T21	Two concurrent edits	One revision wins; stale writer gets 409
T22	Two visitor accounts and guessed record IDs	No cross-workspace reads, writes, exports or model runs
T23	Direct Data API calls attempting to forge assessment/Pc	Denied by grants/RLS
T24	Function termination midway and retry	No partial valid result, no duplicate financial benefit or action
T25	A-B improvement creates A-C conflict	Candidate one blocked; all pair evidence visible
T26	Catalogue/candidate incomplete or unsupported	No candidate pass
T27	Repeated idempotency key	Same payload returns same result; changed payload gets 409
T28	No contact, lost acknowledgement, duplicate packet, stale candidate	Correct simulated state, no fake execution
T29	Reference INR economics	100000 before, 1000 after, 99000 reduction, 79000 net
T30	Mixed currency, missing loss, unsupported Pc, negative net	No unqualified aggregate or fabricated positive savings
T31	Footprint holes, disconnected pieces, boundary points, repeated asset	Correct inclusion and no double count
T32	Invalid/datetime polygon and missing exposure coverage	Reject unsupported geometry or display explicit coverage gap
T33	Missing vulnerability assumption	Exposure available; damage unavailable/partial
T34	Prompt injection in imported text	No tool privilege expansion or policy mutation
T35	Invented finding ID, omitted required finding, wrong template	Reject/fix host packet; no accepted false citation
T36	Groq 429, timeout, absent key, refusal or malformed schema	Deterministic fallback clearly labeled
T37	New assessment during model inference	Old response marked stale
T38	Cold start, browser refresh and sign in on another device	Persistent account sees saved data; no local filesystem dependency
T39	Fresh public browser and auth callback	Live app works; real model trace verified with owner's key
T40	Keyboard and 390/768/1440 px screens	Core workflow accessible, no clipped actions or color-only meaning

Agent evaluation

Create 50 frozen prompts: 20 ordinary evidence investigations, 10 missing-data requests, 10 conflicts/changed versions and 10 adversarial/unsupported cases. Label expected permissible tools/templates before evaluation. Report correct selection count, source fidelity, fallback count, latency and token cost. A template-only run cannot pass the live Groq integration requirement. Schema adherence is not proof of useful tool choice.

Performance protocol

On the actual public tier, measure ten cold-start samples when practical and at least thirty warm repetitions per workload. Warm API target: queue read p95 ≤ 2 seconds for 500 cases, ordinary single encounter audit ≤ 2 seconds, default three-object comparison ≤ 5 seconds and complete normal agent packet ≤ 30 seconds. These are proposed targets, not measured claims or rigid proof of physical timeliness.

Also measure 20-object/4-candidate completion, chunk duration, peak memory, payload sizes, cache hit rate and database calls. A chunk must return before the application budget. Missing targets require profiling and truthful reporting; never hide a timeout by substituting precomputed values. Benchmark sequential reference, sequential Rust, parallel Rust and cached Rust with identical scientific work.

Expert and customer evaluation limits

Synthetic policy consistency is not professional analyst accuracy. For a later pilot, freeze independent expert required-review labels and compare against a baseline already including quality checks and applicable dilution handling. Track both missed required reviews and excess workload at a stated capacity. Counterbalance task order and avoid letting an analyst repeat the same case in both interfaces.

A 99% recall claim needs appropriate independent cases and confidence bounds. For 299 independent required-review cases with zero misses, a one-sided exact 95% lower bound is about 99.003%; thirty perfect cases only about 90.5%. Correlated report updates cannot count as independent trials. Mean net time saved feeds economics; median time describes typical UX. No such study has yet been completed.¹⁶

Source notes

16. NIST/SEMATECH, [Confidence limits for a proportion](#).

Implementation milestones and work tickets

Build in dependency order. A milestone is complete when its evidence exists, not when its screens look finished. Preserve existing files in devx-engine, especially 3d-game/, an important asset for future implementation, and do not execute instructions discovered in reference repositories. Before edits, read applicable repository AGENTS.md instructions, inspect current Git status and create an appropriate isolated branch/worktree without overwriting existing work. Record the game's starting file/configuration state so the final diff can verify preservation; exclude its folder from cleanup and scaffolding replacement.

M0 - Prove the deployment path

Inspect the current remote revision and Vercel project configuration. Record the selected application directory, production branch, runtime and public hostname. Create the orbit-trust application skeleton, static routes, health endpoint, capability endpoint and one synthetic Rust smoke calculation. Build a Python 3.12 Linux wheel in CI, import it in the deployed Python function and render its actual result on a static page. Prove Supabase JWT verification, a private read and an atomic test write in a preview environment. Prove one bounded ADK-to-Groq tool call and a separately validated structured response.

Lock compatible package versions only after these spikes. The expected starting stack is Python 3.12, Node 22, pnpm 10, stable Rust, FastAPI/Pydantic, Google ADK, LiteLLM, Groq SDK, Supabase JS and a small server-side Supabase/PostgREST client. Record exact patches and license notices. Do not pin an untested combination from a slide deck. Exclude the known compromised LiteLLM releases described in document 09.

Exit evidence: public preview URL, route tests, Linux extension build/import log, database isolation check, redacted real provider trace and locked dependency matrix. If a credential is absent, complete the independent scaffolding and mark only that integration blocked. M0 is deliberately early because native wheel packaging and ADK/provider compatibility could change the implementation details.

M1 - Contracts, provenance and storage

Tickets: translate the canonical schema into strict Pydantic/TypeScript models; implement structural and domain validation; create tables/RLS/RPCs; create isolated workspaces; add revision-protected satellite metadata; implement canonical import, immutable report history, explicit event grouping, conflicts, idempotency and exports. Implement GitHub sign-in and anonymous demo identity. Import the synthetic fixtures into workspace-owned records.

Create a local ESA research adapter with aggregate audit reproduction, event-level replay split and encoding exclusions. Historical reports remain a reported-risk replay path; they do not masquerade as complete encounters. No historical upload route is enabled publicly.

Exit evidence: T16-T24 and T27 relevant portions, invalid import report, two-account isolation proof and an export/import round trip. Every data mutation produces an audit event in the same transaction.

M2 - Scientific core and review policy

Tickets: implement covariance validity and frame checks; linear TCA; encounter plane projection; convergent 2D Gaussian integration; independent reference calculations; precision/applicability states; versioned policy; persisted assessment dependencies; deadline clock; urgent acknowledgement floor and reopening. Add reference invariance, failure and threshold tests before presenting risk values in the UI.

Exit evidence: T01-T15, independent reference report and measured single-encounter latency. Document which unsupported domains remain unsupported. No LLM is needed to pass this milestone.

M3 - Fleet and response simulation

Tickets: create candidate catalogue inputs and revisions; enumerate all unordered pairs over the supplied interval; compare baseline and supplied alternatives; implement bounded chunks, DB leases and safe continuation; preserve affected and unchanged pair results; derive candidate dispositions; add contact-window and packet state simulation. Write sequential and parallel benchmark modes with cache controls.

Exit evidence: T24-T28; A-B improvement that creates A-C concern is blocked; interrupted work resumes without publishing a partial pass; 20-object limits are enforced. Command simulation has no network connector to a satellite or ground station.

M4 - Consequence models and reentry sandbox

Tickets: exact decimal monetary inputs; conditional spacecraft loss breakdown; expected-loss comparison; uncertainty scenarios; regional polygon validation; point inclusion with holes/boundaries; exposure deduplication; missing-data coverage; optional conditional damage calculation. Reuse TerraSight's separation of observations and decisions, not its rover polygon assembler or scoring equations.

Exit evidence: T29-T33, fixture totals reproduced and all economic assumptions visible. Reentry inputs cannot alter orbital policy. A high asset value cannot lower a review priority.

M5 - Bounded agent investigation

Tickets: ADK workflow; allowlisted version-bound tools; shared token/time admission; investigator tool choice; tool-free formatter; host validation; template rendering; required-finding inclusion; source/provenance trace; cancellation and stale-result handling; deterministic fallback. Freeze the 50-item evaluation set before measuring.

Exit evidence: T34-T37, one real Groq run in the public app, evaluation counts and token/latency report. A valid JSON shape is not the whole evaluation. The investigator must select useful evidence within its budget.

M6 - Complete mission console

Tickets: reusable accessible sidebar/panels; server-paginated queue; case evidence/timeline; fleet matrix; simulated response status; satellite add/edit; economics; separate reentry tab; activity feed; settings, data notices and export. Include skeleton/loading, no-data, failed, unsupported, stale, quota-exhausted and offline states. Expose deterministic results before the model response arrives.

Exit evidence: T19-T21 and T38-T40, screenshots at three widths, keyboard review, fresh-session persistence and zero inaccessible urgent cases. Test browser back/forward and fixed routes with query IDs on the deployed URL.

M7 - Release, evidence and submission

Run the applicable suite once against the release candidate; address failures, then rerun affected checks. Measure public latency and bundle size. Review secrets, data notices and license status. Verify that 3d-game/ and its deployment settings remain intact; if shared configuration changed, check that the game build/deployment behavior is preserved. Record all requirement outcomes in document 17, populate BUILD_STATUS, rehearse the three-minute demonstration and make the recording.

Publish the authorized final release through the existing Vercel workflow after preview verification. A coding request alone does not override a separate environment approval requirement; ask only when an actual required permission is missing. Do not imply that this document has already authorized sending customer outreach or incurring paid charges.

Exit evidence: stable public URL, release commit, test/benchmark results, recording, pitch text, reproducible local instructions and explicit limitations. A local-only build fails the public submission requirement.

Cut order if implementation constraints emerge

Keep the scientific audit, core queue, three-object candidate reversal, genuine bounded ADK/Groq run, persistence and public URL. Reduce optional animation, custom map tiles, dashboard personalization, large benchmark sizes and report styling first. The separately scoped reentry and finance requirements remain required unless the owner explicitly accepts a scope change. Never silently replace Rust with displayed fixture constants or the live agent with a fabricated trace.

Ticket completion template

Each ticket records: requirement IDs; input/output contract; files changed; dependency; acceptance test IDs; measured result; outstanding assumptions; source/method versions; reviewer notes. Use GitHub issues only if the owner requests external issue creation. Otherwise keep the backlog in docs/BUILD_STATUS.md.

Deployment and operations runbook

Deployment target and configuration gate

The existing repository is <https://github.com/Adhirajsingh2507/devx-engine>. It contains toolkit resources and the important 3d-game/ app reserved for future implementation. Preserve the game's files, assets, configuration and any existing deployment. Use orbit-trust/ as the intended Root Directory for the ORBIT-TRUST deployment only. The existing project's actual ID, linked branch, preset, Root Directory, deployment protection and environment variables have not been inspected through the owner's account. Record and verify them before changing a live project. If that project serves the game, keep it intact and resolve a distinct ORBIT-TRUST deployment target with the owner before changing settings.

Keep one Vercel project for the ORBIT-TRUST frontend/backend composition; this does not authorize consolidating or replacing a game deployment. Next.js exports the fixed page list to static files; root app.py exposes FastAPI; exported assets are copied into public/ for CDN delivery. Vercel documents FastAPI detection and a single Python function, with public/ used for static assets. The build must demonstrate correct precedence for /api/v1/* and the fixed web routes. A blanket HTML fallback for API failures is forbidden.¹²

Use a reproducible build script with a frozen pnpm lockfile. It builds apps/web, copies the export into the app root's public directory and verifies every expected route. Do not set the Vercel preset to Next.js merely because the UI uses Next.js: the selected project entry is the FastAPI/static composition. Verify the current framework preset, build command and output handling in M0 instead of inventing a known-working vercel.json. Keep the successful configuration in deploy/ with the tested Vercel CLI version.

Native Rust packaging

Build the PyO3 extension with a pinned maturin toolchain in Linux CI for the exact Python/runtime architecture selected in M0. Prefer a compatible manylinux wheel with no unbundled external libraries. Name the artifact by package version; record wheel SHA-256, source-tree digest, Rust version, target, Python ABI and dependency lock hashes in vendor/wheels/build-manifest.json.

The deployable commit must already contain the tested wheel and manifest before Vercel's automatic deployment runs. A race in which CI builds the wheel after the same commit has auto-deployed is not acceptable. Prepare the wheel on the implementation branch, commit it with the matching source, verify it in CI, and only then promote that commit. A subsequent source change requires a new matching artifact. Keep a strict small-artifact limit and use a versioned authenticated artifact store later if native bundles become large.

Install the wheel during Python dependency installation. Do not rely on compiling Rust in an unspecified build stage after Vercel has already installed Python requirements. Verify extension import and an independent reference case on Linux and the actual preview runtime. Local Windows development uses a Windows-compatible build, never the Linux wheel. Pin releases and verify hashes; do not download arbitrary binaries from a user-supplied URL.^{13, 20}

Environment variables

Name	Scope	Meaning
NEXT_PUBLIC_SUPABASE_URL	Browser/build	Supabase project URL
NEXT_PUBLIC_SUPABASE_PUBLISHABLE_KEY	Browser/build	Public client key; grants/RLS still enforce access
SUPABASE_URL	Server	Same project endpoint
SUPABASE_SERVICE_ROLE_KEY	Server secret	Restricted use for validated transactional RPC writes; never bundle

Name	Scope	Meaning
GROQ_API_KEY	Server secret	Owner's free-tier key
GROQ_MODEL	Server	Raw provider model ID; default openai/gpt-oss-20b if available
ADK_LITELLM_MODEL	Server	Adapter ID groq/openai/gpt-oss-20b
APP_PUBLIC_ORIGIN	Server	Exact approved public origin
APP_MODE	Server	public_synthetic for submission
POLICY_VERSION	Server	demo-2.0
APP_REQUEST_BUDGET_SECONDS	Server	45; persistence reserve is 5
ORBIT_CORE_THREADS	Server	1 initially; change only with measurement
GROQ_ENABLED	Server	Explicit integration switch, false yields labeled fallback

Store only placeholders in `.env.example`. Use separate local, preview and production values. Do not print credentials in logs, health checks, traces or generated reports. Supabase may expose newer server secret key names; the variable above is an application convention, not a requirement to use a legacy key type. Map a supported server credential deliberately and keep its use on the backend.

Supabase setup and login

Create/select the owner's free project, apply migrations, verify grants/RLS and seed only original synthetic templates. Enable anonymous sign-ins for immediate isolated trial access. Enable GitHub OAuth for persistent accounts; configure a GitHub OAuth application's callback to the exact Supabase Auth callback URL shown in its dashboard. Keep that OAuth client secret in the provider configuration, not frontend code. Add the final public app callback and explicit local/preview callbacks to Supabase's allowed redirect list. ^{17, 27}

The statically exported `/auth/callback/` page uses the browser Supabase client with a tested PKCE flow and exchanges the returned code. It is a client page, not a Next.js server route handler. Test refresh, errors, sign-out, another browser and the anonymous-to-persistent transition. Prevent arbitrary external return URLs; allow only known local paths.

Anonymous identity is still a Supabase authenticated user. A lost browser session cannot be treated as a recoverable persistent account. Offer explicit export before sign-out. In version 2.0, signing into an existing GitHub account opens that account's workspace; it does not silently merge the anonymous one. Offer validated export/import to move synthetic case inputs, with fresh runtime IDs and preserved source lineage. This avoids ambiguous account linking while satisfying persistent accounts. ¹⁷

Authentication, JWT expiry, quotas, leases and retention use actual UTC wall time. The fictional demo clock only affects scenario calculations and review deadlines.

Runtime and quota controls

Set a verified Vercel function duration above the 45-second application budget; 60 seconds is the initial configuration target if supported by the resolved endpoint/plan. Stop expensive work before the persistence reserve. Check current Python bundle and response limits at build time, and keep archives/training dependencies outside the deployment. No paid service is provisioned automatically. ¹³

Use the database for one active operation per workspace, one global live Groq run, admission counts and expiry. Distinguish application runs from provider requests: allow at most three investigate runs per minute per user, ten per UTC day for anonymous users and fifty for persistent users; each run has the tighter provider-call/token limits in document 09. These are additional caps, not a guarantee of provider capacity. Enforce a shared daily provider token reservation below the actual account allowance; release unused

reservation after accounting. Retry at most once within the remaining run budget, respect Retry-After and otherwise show fallback.⁹

Add signup abuse controls supported by the chosen Supabase plan, such as CAPTCHA, and cap workspace creation to one active demo workspace per identity. Anonymous churn means user limits alone cannot protect the key; global provider admission and fail-closed quotas are required. Never use recurring artificial traffic to evade free-project pause behavior. Show a useful unavailable message and restoration runbook instead.¹⁹

Vercel Hobby restricts use to personal, non-commercial use. Verify that the actual hackathon deployment qualifies; a later paid commercial pilot needs an eligible hosting plan and operating budget. This does not establish that a commercial product can operate indefinitely at zero cost.²⁶

Release and recovery

1. Create a preview from the deployable commit. Apply backward-compatible migrations to the appropriate database environment; never reset production to run tests.
2. Verify extension import, authentication, two-user isolation, metadata persistence, default encounter, fleet reversal, finance, reentry, real Groq trace and API/static routing.
3. Record cold/warm results and quotas. Confirm the public judge URL does not require Vercel team login. Restrict preview secrets and avoid making private test cases public.
4. Promote the verified commit through the existing deployment workflow; record the production URL and release ID. Smoke test from a fresh browser.
5. If release fails, roll back application deployment to a known compatible commit. Do not undo additive migrations by dropping data. Restore from a tested export/backup only with an explicit recovery procedure.

Export a synthetic demo backup before judging. Free-plan backup features vary; confirm actual facilities, and keep a manual versioned JSON export sufficient for the demo. Redact tokens and personal data from incident records. Retain anonymous inactive workspaces for 7 days and agent trace metadata for 30 days as proposed app policy; implement deletion through a documented owner-run maintenance action initially, not a hidden continuously running worker. Persistent accounts retain case data until owner deletion, with the stated public per-workspace limits. Implement delete-workspace/account UI only with explicit confirmation of that concrete destructive action.

Deployment failures to diagnose

Wheel import error: inspect ABI/architecture/library dependencies and manifest; do not return fake values. Static route 404: inspect export tree and route precedence. OAuth loop: compare exact origins/callbacks and PKCE storage. Database 403: inspect grants, membership and JWT, not a blanket RLS disable. Groq 429: inspect account limits and admission counters. Function timeout: inspect bounded chunk size and response checkpoint, not a background thread. Existing game replaced unintentionally: restore the prior Vercel Root Directory and select the intended project with the owner.

Source notes

9. Groq, [Rate limits](#).

12. Vercel, [FastAPI deployment](#), refreshed 11 September 2026.

13. Vercel, [Python runtime](#) and [Function limitations](#).

17. Supabase, [Anonymous sign-ins](#).

19. Supabase, [Billing on Supabase](#).

20. PyO3, [Guide](#); Rayon, [API documentation](#).

26. Vercel, [Hobby plan](#), refreshed 11 September 2026.

27. Supabase, [Login with GitHub](#), refreshed 11 September 2026.

Business model, validation and hackathon presentation

Positioning

ORBIT-TRUST is an evidence and response-review workspace for teams that already receive conjunction warnings. Its proposed advantage is the connection between changing evidence, review deadlines, supplied fleet response alternatives and a reproducible decision record. It is not a new tracking network or a promise to eliminate orbital collisions.

Start with a flight-dynamics team at a small LEO operator or an outsourced mission-operations provider. The daily user reviews conjunction cases; the operations lead approves workflow purchases; security and flight-dynamics specialists gate operational adoption. A university CubeSat team may be accessible for usability feedback but does not automatically represent a paying commercial buyer.

Existing alternatives and differentiation test

Alternative	Evidence available	Consequence for positioning
Provider alerts plus spreadsheets and internal scripts	A plausible workflow baseline to verify in interviews	Compare against actual review practice, not a deliberately weak Pc-only table
Kayhan	Public documentation discusses risk metrics and dilution ⁶	Uncertainty-aware metrics alone are not novel
Omitron OwlEye	Markets conjunction/risk assessment and operational support ²¹	Do not claim the first decision-support product
Neuraspace STM OPS	Markets operational space-traffic tools ²²	Workflow automation and maneuver support already have competitors
ESA CREAM	Research into collision-avoidance automation ²³	Automation itself is not sufficient differentiation

Public marketing establishes advertised capabilities, not completeness, efficacy or a verified feature absence. ORBIT-TRUST's gap hypothesis is easier reconciliation and auditability for an underserved workflow. It needs side-by-side analyst testing. An eventual advantage could be operator-approved adapters, expert-labeled evidence conflicts and integration history; the demo does not yet possess that advantage.

Proposed commercial model

Sell a team workspace subscription with a defined reviewed-event allowance and support level. A later enterprise tier can add private deployment, provider integrations, access controls and audit retention. Charge for the workflow and support, not a percentage of a claimed satellite saved. Avoid transaction-linked financial promises that the evidence cannot establish.

Initial pricing is an interview hypothesis: a time-limited paid pilot at INR 25,000 per month, followed by a small-team plan around INR 50,000 per month if verified value supports it. These are original proposal numbers, not observed market prices or accepted willingness to pay. Quote taxes, data-provider costs and private-hosting costs separately if applicable. Do not publish a checkout before validating plan eligibility, terms and support costs.

For a labor-value example, 600 reviewed cases/month with a measured mean net reduction of 8 minutes/case produces 80 hours/month. At an explicitly assumed loaded analyst cost of INR 1,500/hour, the modeled labor capacity value is INR 120,000/month. A hypothetical INR 50,000 subscription leaves INR 70,000/month before integration, training and other costs. The case count, time saving and labor rate are assumptions; saved capacity is not automatically a cash payroll reduction.

Contribution per customer equals revenue minus compute/inference, database, licensed data, support and amortized onboarding costs. Keep customer acquisition and general R&D separate. Free hackathon quotas do not prove sustainable margins. No defensible TAM or market-share estimate is currently available; a later bottom-up model should count independently verified qualified operators, realistic annual contract values and reachable conversion rates instead of multiplying every satellite by a price.

Unit-economics hypotheses

These scenarios are arithmetic illustrations for discovery, not supplier quotes, forecasts or measured business performance. Replace every cost with actual pilot data before investment or pricing decisions.

Monthly per-customer input or derived metric	Base hypothesis	Stress hypothesis
Subscription revenue / ARPU	INR 50,000	INR 50,000
Compute, inference and database allocation	INR 5,000	INR 10,000
Licensed-data allocation	INR 5,000	INR 10,000
Support and amortized onboarding	INR 15,000	INR 20,000
Total cost to serve	INR 25,000	INR 40,000
Gross contribution / margin	INR 25,000 / 50%	INR 10,000 / 20%
Acquisition cost per customer, assumed	INR 100,000	INR 200,000
Acquisition payback	4 months	20 months
Retained lifetime, assumed	24 months	12 months
Simplified contribution LTV	INR 600,000	INR 120,000
LTV / acquisition cost	6.0	0.6

Payback equals acquisition cost divided by monthly gross contribution. Simplified LTV equals monthly gross contribution times assumed retained lifetime; it omits discounting, expansion and changing support burden. At a hypothetical INR 200,000/month fixed operating expense, the base contribution requires eight paying customers just to cover that expense, before acquisition spending. The stress case is unattractive: payback exceeds assumed lifetime. These examples make the support/data-cost risk visible instead of equating free inference with a high-margin business.

The first distribution step is a small set of founder-led operator/service-provider conversations, followed by one to three scoped pilots. Ten customers require repeatable onboarding and a verified budget owner. Reaching a hundred would likely require service-provider distribution and reusable data adapters. A thousand qualified buyers is not assumed in this narrow market. Reentry exposure has a different operational workflow; treat it as an experimental extension until a buyer validates its value, rather than charging the conjunction buyer for every demo feature.

Customer discovery plan

Conduct 8-12 structured conversations spanning operators and service providers, with permission and without asking for sensitive conjunction data. Ask how they handled their most recent difficult alert; which evidence required clarification; who made the review decision; which system already handles that task; and what integration/security work a pilot would require. Show the three-object reversal only after learning their current process. Do not ask whether they simply like the idea.

For a pilot, agree on expert labels, representative cases, the existing-tool baseline and a stop condition before testing. Measure mean review time, missed required reviews, excess escalations and time to retrieve supporting evidence. Seek written interest in a specific evaluated workflow and a named operational sponsor.

No interviews, letters of intent, revenue, contracts or partnerships are claimed in this handoff.

Three-minute live demonstration

Time	Action	Message
0:00-0:25	Open queue with source updates and urgent count	A lower reported probability can coexist with unresolved evidence and a review deadline
0:25-0:55	Inspect case timeline and deterministic findings	The system preserves provenance and explains why this case needs attention
0:55-1:20	Run ADK investigation; show actual tool trace	The agent selects evidence and drafts a bounded request; policy remains deterministic
1:20-1:55	Compare baseline and two supplied response candidates	Candidate one helps A-B but creates A-C concern; fleet coverage changes the disposition
1:55-2:15	Show simulated communication and separate financial example	Acknowledgement is not execution; monetary benefit is conditional expected loss
2:15-2:40	Open separate reentry sandbox and change footprint	Exposure changes under a supplied footprint; this is not a predicted crash location
2:40-3:00	Export packet and show measured test results	Every accepted fact is tied to its inputs, assumptions and method

Keep the INR 79,000 financial example explicitly labeled as a separate economic reference scenario unless the implementation binds it to matching Pc inputs. Do not accidentally associate those probabilities with the fleet fixture's different probabilities. Begin with a pre-seeded workspace; calculations execute on demand. If the provider fails during presentation, show the honest fallback and previously recorded real-run evidence with its date, rather than inventing live agent activity.

Suggested pitch

“Satellite teams do not just need another warning. They need a defensible review before the deadline. ORBIT-TRUST audits the evidence, keeps urgent uncertainty visible and tests supplied response options against the loaded fleet. In our synthetic example, a response that improves one conjunction creates another. Deterministic calculations catch that change; bounded agents help assemble the evidence request. The separate reentry sandbox reports exposure under stated footprints. We are testing whether this workflow reduces analyst effort without missing required reviews.”

Submission evidence and judging

The supplied Google ADK PDF explains components; it is not a verified competition rulebook or scoring rubric. Hackathon name, deadline, mandatory services, allowed pre-existing code and submission fields remain unconfirmed. Map the final submission to the actual rules when they are available. Do not invent numerical judging weights or claim Google ADK/Groq compatibility guarantees eligibility.

Prepare the public URL, repository URL and release commit, three-minute recording, architecture diagram, source/license notices, reference checks, measured latency table, real provider trace and clear reuse disclosure. Describe the existing projects as prior work and identify the new domain-specific implementation. Do not claim all components were built during the hackathon if they were adapted from existing repositories.

Judge questions with defensible answers

Why agents? They choose among bounded evidence tools and organize an information request; deterministic services calculate and enforce policy. Why Rust? It provides a shared tested numerical batch core; performance benefit will be reported from measurements, not assumed from the language. Why not existing products? This is a workflow hypothesis requiring validation against incumbents. Can it prevent a collision? The demo supports review and supplied scenario comparison; it cannot certify or execute a real maneuver. Does the reentry map predict casualties? No: it evaluates exposure under a supplied footprint and optional stated damage assumptions.

Winning cannot be guaranteed. The strongest demonstration is a specific, reproducible decision failure that the application catches, a genuine bounded agent contribution, and clear evidence that the deployed implementation works.

Source notes

6. Kayhan Space, [Risk metrics](#).

21. Omitron, [OwlEye](#).

22. Neuraspace, [STM OPS](#).

23. ESA, [CREAM: avoiding collisions through automation](#).

Sources and evidence register

Sources were inspected during the project research and refreshed where deployment choices changed, through 11 September 2026. S-numbers are stable source identifiers used throughout this package. Official pages can change: recheck provider/runtime details during M0. Scientific equations and demo-policy choices are specified separately from provider recommendations.

ID	Publisher and source	Use and limitation
S01	NASA, Spacecraft Conjunction Assessment and Collision Avoidance Best Practices Handbook , 2023; NASA CARA, Close Approach Risk Mitigation	Operational risk review and mitigation context; not an endorsement of the demo policy
S02	ESA, Reentry background	Reentry is a distinct physical problem; does not supply this demo's footprint
S03	ESA, Re-entry safety	Trajectory, fragment survival and population-risk context
S04	CCSDS, Conjunction Data Message, 508.0-B-1 with corrections	Standard message context; this release accepts a smaller canonical JSON contract and is not a full standards-compliant parser
S05	NASA NTRS, Covariance realism research record	A mathematically valid covariance is not automatically a calibrated error distribution
S06	Kayhan Space, Risk metrics	Risk/dilution context and evidence that incumbent tooling discusses this issue
S07	Google ADK, LiteLLM integration	Adapter selection and dependency security notice; exact package combination still requires integration testing
S08	Groq, Supported models	Candidate provider model availability; actual account access must be verified
S09	Groq, Rate limits	Provider admission constraints, not guaranteed free capacity for every account
S10	Groq, Tool use and Structured outputs	Separate tool investigation and schema-constrained formatting; verify model support
S11	Google ADK, Workflow agents	Orchestration concepts; runtime behavior depends on locked versions
S12	Vercel, FastAPI deployment , refreshed 11 September 2026	Python app detection, static assets and deployment composition; M0 verifies the combined build
S13	Vercel, Python runtime and Function limitations	Packaging/runtime constraints; inspect live limits rather than assuming a persistent server
S14	Next.js, Static exports	Static UI shell and unsupported dynamic server features
S15	ESA Kelvins, Collision Avoidance Challenge data and Rules	Historical replay dataset; aggregate archive audit included, raw data and redistribution rights not bundled
S16	NIST/SEMATECH, Confidence limits for a proportion	Exact-binomial confidence context; correlated report rows are not independent trials
S17	Supabase, Anonymous sign-ins	Anonymous identity, persistence and abuse-control considerations

ID	Publisher and source	Use and limitation
S18	Supabase, Row Level Security	Client authorization model and service-key caveats
S19	Supabase, Billing on Supabase	Plan facilities and limits; verify actual project status and do not assume perpetual free uptime
S20	PyO3, Guide ; Rayon, API documentation	Native Python binding and bounded parallel batch design
S21	Omitron, OwlEye	Advertised competing operational capability; efficacy not independently tested
S22	Neuraspace, STM OPS	Advertised competing workflow; no claim of absent features
S23	ESA, CREAM: avoiding collisions through automation	Existing research direction in automation
S24	Adhirajsingh2507, repository snapshots in document 03	Source-level reuse assessment, not executed application verification
S25	Sitam Meur / ML Kolkata, Google ADK, component by component, 11 September 2026; supplied file google-adk-component-by-component (1).pdf, 20 pages	Supplied component deck, especially pages 4-9 for agents, tools, schema and workflows; bundled for context, not the competition rulebook
S26	Vercel, Hobby plan , refreshed 11 September 2026	Personal/non-commercial restriction and account limits
S27	Supabase, Login with GitHub , refreshed 11 September 2026	OAuth provider/callback setup; adapt examples to a client-rendered static callback

Repository snapshot evidence

- [TerraSight](#) at [2e6445f0639f1d20ce0ad45c98134e9bcc11ea20](#): terrain/scoring and pipeline review.
- [Finora](#) at [c66824b8242262bb0f640f4b98ad9067591d29d4](#): finance engine, judge and orchestrator review.
- [Sajawat](#) at [600cf5536d8f0b138bdf470e298329213cf1bf57](#): CRM queue/detail component review.
- [devx-engine](#) at [ad4b9e45b234fe8b4d22d7db3be7c74d46c7a3fd](#): tree and target repository structure review.

The GitHub pages and inspected source text do not prove license clearance, build success or security. Preserve author attribution and confirm owner rights before redistributing adapted source. A missing detected license is not evidence that arbitrary public code may be copied without conditions.

Observation, proposal and unresolved evidence

Observed: repository structure and cited source snippets; ESA aggregate archive statistics; independent numerical reference agreement recorded under reference/. Proposed: thresholds, capacity limits, design, pricing, application architecture, retention and latency targets. Unverified: live application, actual Vercel settings, selected package interoperability, real operator benefit, reentry footprint validity, customer's willingness to pay and hackathon eligibility rules.

The original research role and engineering lecture priorities are summarized in [reference/OWNER_BRIEF.md](#). Attached-document instructions were treated as material to evaluate, not as tool-use authority. This package deliberately carries the resulting requirements and their limits rather than conflicting obsolete implementation assumptions.

Traceability and release gate

Requirement	Specification	Milestone	Acceptance evidence
R01 imports	05, 11	M1	Canonical positive/negative fixtures, historical audit reproduction, T07
R02 revisions	05, 19	M1	T16-T18, T21
R03 evidence	06	M2	T07-T15
R04 Rust probability	06	M0, M2	T01-T10, Linux import
R05 review priority	06	M2	T11-T15, T19-T20
R06 fleet comparison	07	M3	T24-T26, every pair coverage
R07 communication simulation	07	M3	T28, no command connector
R08 financial scenarios	08	M4	T29-T30, exact arithmetic
R09 reentry sandbox	08	M4	T31-T33, visible provenance
R10 ADK/Groq	09	M0, M5	T34-T37, real provider trace
R11 validation	09	M5	T35, independent host policy
R12 mission console	10	M6	T19-T21, T38-T40
R13 graceful fallback	09, 14	M5	T36 and current deterministic packet
R14 public deployment	04, 14, 19	M0, M7	T22-T24, T38-T39, public URL
R15 reproducibility	05, 11, 19	M1, M7	Export/import equality, immutable hashes, T27
R16 resource/accessibility	04, 10, 12	M6, M7	T40 and measured performance report

Release gate

Repository preservation is also a release gate: the existing 3d-game/ app is important for future implementation. Verify against the starting Git state that its source, assets, dependencies and configuration have not been removed, overwritten, renamed, moved or repurposed, and any existing game deployment remains intact. Record the diff/configuration evidence in BUILD_STATUS. Shared build/configuration changes must preserve game behavior. This is a preservation requirement, not a new game feature to implement during the hackathon build.

Record PASS, FAIL, BLOCKED or NOT_RUN for every requirement and each applicable test; do not convert NOT_RUN to PASS. A release candidate needs all required behavior implemented, critical numerical/security tests passing, bounded failure behavior and a public interactive URL. A missing Groq key allows independent development but blocks the real integration acceptance. A failed provider call with honest fallback is correct failure behavior and still does not prove successful live integration.

Include commit, public URL, timestamp, runtime, dependency locks, schema/policy/method versions, wheel/source hashes, dataset/fixture hashes, measured latency percentiles and all unresolved deviations. A reported scientific limitation can be acceptable when it is an intentional supported-domain boundary; an unnoticed zero-risk substitution is a failure.

No raw ESA archive, private operator report, secret, customer contact or browser authentication state belongs in the public build or evidence bundle. No customer outreach, paid purchase or external publication of private

data is part of the implementation plan.

Build status starter

Project: ORBIT-TRUST 2.0. Target: devx-engine/orbit-trust. Hosting: Vercel. Persistence: Supabase. Provider: Groq. Current state at handoff: specification prepared; application implementation NOT_STARTED. Application test outcomes: NOT_RUN. Public URL: owner account configuration pending. Provider integration: NOT_RUN. Operator validation and commercial traction: NONE CLAIMED.

The implementing chat should copy this starter into docs/BUILD_STATUS.md and update it as evidence arrives. Keep a decisions/deviations table with date, reason, requirement impact, owner approval if actually required and replacement acceptance evidence. End each work session with the next executable ticket and any credential gate, so later sessions resume instead of rebuilding the plan.

Glossary and implementation recovery

Term	Meaning in this project
Conjunction	Predicted close approach between two objects; not proof of impact
TCA	Time of closest approach within the specified model and interval
Pc	Model-based probability under stated geometry/error assumptions
Covariance	Error-distribution model; mathematical validity and empirical calibration differ
Dilution	Greater modeled uncertainty can lower computed Pc in some regimes
Hard-body radius	Sum of the supplied object collision radii in the demo model
Evidence state	Whether inputs support the declared assessment; separate from urgency
Review deadline	Supplied latest action timing minus review allowance, not automatically TCA
Candidate	Supplied trajectory/planning alternative; no burn is derived by the demo
Coverage	Objects, pairs, intervals and evidence actually evaluated
Exposure	Population/asset inventory intersecting a supplied footprint
Expected loss	Probability-weighted conditional monetary outcome under stated assumptions
Provenance	Source, revisions, hashes and methods behind a fact
Lease	Time-limited database ownership of a bounded operation
Deterministic fallback	Host-generated packet from checked facts when the model cannot run

First actions in a new account

Read `START_HERE` and the bootstrap prompt. Inspect repository access and the live target tree; compare with the recorded snapshot. Follow M0 before building all screens. The ZIP is the most complete transfer. If only the combined Markdown is supplied, reconstruct schemas/ and fixtures/ from the named JSON appendices before running validation. The original ADK PDF is an optional contextual reference; its component summary is included in the text pack.

Access to GitHub is not access to Vercel, Supabase or Groq secrets. The implementer can complete source code, local fixtures and tests without the owner's cloud credentials, then list exact remaining account actions. Never ask the owner to paste keys into a public repository or report. If GitHub is temporarily unavailable, use the documented interfaces and defer source reuse review; do not pretend the source was freshly inspected.

Handling conflicts

If a newer library removes a model/adaptor, perform a small replacement compatibility test, update the decision log and preserve the tool/validation contracts. If Vercel routing changes, resolve the static/FastAPI composition in M0. If native packaging is blocked, retain the Rust requirement and surface the packaging gate; a temporary local reference mode must identify itself. If actual hackathon rules prohibit pre-existing components or require a different deployment, present the concrete rule and minimal affected change to the owner.

If the owner later asks for actual orbital propagation, autonomous command execution, real reentry prediction or operator data ingestion, treat it as a new engineering scope. It requires new scientific models, input agreements and acceptance criteria. Do not enable it by removing the simulation label.

Interpretation of the original ideas

Terrain reuse becomes a separate exposure tool because rover path planning cannot steer debris. Finance reuse becomes explicit loss mathematics because personal budgeting rules do not model spacecraft losses. CRM reuse becomes an evidence/timeline workspace because sales stages do not represent review decisions. ADAS becomes a bounded response simulation because the demo lacks onboard authority and qualified flight software. Multiagent design becomes a few useful roles around deterministic services because multiple agreeing models do not prove correctness. Rust is an implementation choice to benchmark, not a guarantee of faster communication or real-time protection.

These corrections preserve the useful intent of the suggestions while making each demonstration claim testable.

Supabase database, authorization and transaction specification

Identity and permissions

Every private record belongs to a workspace. Membership roles are owner, analyst and viewer. Version 2.0 creates one owner membership and one active demo workspace per authenticated identity; team invitation is deferred. An analyst may import, edit metadata and record reviews; a viewer may read/export; only an owner may reset/delete the workspace or change settings. Anonymous demo owners have the same within-workspace operations but lower quotas and limited retention.

Validate Supabase JWT signatures against the project's supported signing configuration. Check issuer, audience, expiry and authenticated subject; cache JWKS with bounded refresh and fail closed on unknown/invalid keys. Use the configured server verification method for legacy signing configurations rather than decoding an unsigned token. Identity comes from the verified token, never an actor_id supplied in JSON.

Client reads use the user JWT with row-level security. All domain writes go through FastAPI, which verifies identity and membership, then invokes narrow server-only transactional functions. A service credential bypasses normal RLS; each function must independently enforce the verified actor's membership and permitted action. Revoke function execution from PUBLIC, anon and authenticated roles unless explicitly designated as a safe read helper. Never expose a generic arbitrary-SQL RPC.¹⁸

Relational inventory

Use UUID primary keys, timestamptz for UTC times, bigint revision/sequence counters, explicit status CHECK constraints and jsonb for bounded validated snapshots. Use numeric or decimal-string fields for money, not float columns. Every child relationship uses a composite workspace_id + entity_id foreign key where practical to prevent cross-workspace references even inside a server bug.

Table	Keys and important fields
workspaces	id, owner_id, mode, data_revision, queue_revision, next_audit_sequence, demo_clock, created_at, last_active_at
workspace_members	workspace_id + user_id unique, role
satellites	workspace_id + id, current_revision, archived flag
satellite_revisions	workspace_id + satellite_id + revision unique, immutable metadata, author, created_at
cases	workspace_id + id, event_key unique per workspace, ordered pair identity, current_revision, workflow_state, assigned_member, latest_assessment_id, urgent_floor
reports	workspace_id + id, case_id, source_name, source_revision, raw_digest, normalized_digest, source_created_at, received_at, raw_json, normalized_json
source_conflicts	workspace_id, case_id, competing report IDs, state and review action reference
assessments	workspace_id + id, case_id, input_revision, input_digest, method/policy versions, evidence_state, proposed_urgency, effective_urgency, deadline, immutable facts/findings, method_status
candidate_sets	workspace_id + id + revision, catalogue revision, interval, canonical snapshot and digest
comparisons	workspace_id + id, candidate_set/revision, input_digest, status, complete pair result snapshot and coverage

Table	Keys and important fields
economic_scenarios	workspace_id + id + revision, probability references, decimal inputs, currency, immutable result/dependencies
reentry_scenarios	workspace_id + id + revision, footprint/layer snapshots, assumptions, immutable result/dependencies
operations	workspace_id + id, kind, input_revision/digest, state, lease_token, lease_expires_at, cursor, progress, result_id and error code
operation_items	workspace_id + operation_id + item_key unique, input_digest, status, result snapshot
agent_runs	workspace_id + id, assessment_id, input_digest, model/prompt versions, status, bounded trace, usage, validation and packet selection
review_actions	workspace_id + id, case_id, actor_id, acknowledged_assessment_id, expected_revision, action, rationale, created_at
communication_packets	workspace_id + id, candidate/comparison bindings, simulation state, contact assumptions and acknowledgement/execution event references
audit_events	workspace_id + sequence unique, actor_id, entity/revision, event type, redacted summary and created_at
idempotency_records	workspace_id + actor_id + route + key unique, payload digest, status, result reference, expires_at
quota_buckets	scope + subject + wall-time window unique, reserved/consumed counters and expiry; server-only
global_leases	named capacity slot unique, holder token, expiry; server-only
public_seed_templates	seed_id + version, immutable synthetic payload/digest; read-only public metadata

Report uniqueness is workspace + case + source_name + source_revision + raw_digest. A different digest with the same source revision is permitted as a retained conflict; a uniqueness rule on source revision alone would incorrectly discard that evidence. Ordered pair identity is canonicalized for grouping while primary/secondary roles remain preserved in each report.

Index queue reads on workspace/effective urgency/deadline/case ID; reports on workspace/case/source/created time; operations on workspace/status/lease expiry; audit on workspace/sequence; assessments on workspace/input digest/method/policy. Query counts over the entire filtered dataset, not the returned page. Store immutable derived results and atomically update current pointers only when their dependencies remain current.

RLS and grants matrix

Object	Browser read	Browser direct write	Server transaction
Workspace/member rows	Own membership permits scoped read	Denied	Validated creation/settings only
User data and derived tables	Membership-scoped	Denied	Role plus dependency checks
Audit/review history	Membership-scoped	Denied	Append through domain operation

Object	Browser read	Browser direct write	Server transaction
Quota/global lease/idempotency internals	Denied	Denied	Narrow server-only operations
Public synthetic template metadata	Read permitted	Denied	Deployment/maintenance only

Enable RLS on exposed user tables and revoke unnecessary table grants. A protected membership helper may use SECURITY DEFINER to avoid policy recursion only with a fixed safe search_path, qualified identifiers, no dynamic SQL, least privileges and explicit execute grants. Test the helper with users having no membership. Do not use user-editable profile metadata to grant roles.

Unauthenticated requests cannot create private records. Anonymous sign-in obtains a real authenticated Supabase identity before POST /workspaces/demo. Application admin privileges never derive from is_anonymous=false alone. Return 404 for inaccessible record IDs to avoid existence disclosure; use 403 when a known workspace member lacks an operation permission.

Required atomic operations

- 1. Create/seed workspace:** verify identity, reserve allowed workspace slot, create membership, copy seed inputs to fresh IDs, write initial revision/audit, return result under an idempotency record. Concurrent duplicate creation must not create two active workspaces.
- 2. Import batch:** lock workspace/cases in stable order, enforce caps and expected revision, validate every supplied object first, insert accepted reports/conflicts, invalidate dependencies and advance queue/audit revision. Structural bundle failure rejects the bundle; per-record domain rejection is reported without inserting that record. Accepted records and their import summary commit together.
- 3. Save satellite/action/settings:** lock current record, compare expected_revision, enforce role and referenced assessment freshness, create revision/action, update pointers/floor/workflow and append audit atomically. A stale acknowledgement cannot clear the latest urgent floor.
- 4. Claim operation:** verify membership, idempotency, quotas and input snapshot; acquire one workspace lease and any provider global slot, issue unpredictable lease token and expiry using database wall time. Do not hold a SQL transaction open during model or numerical work.
- 5. Commit chunk:** reacquire row lock, verify lease token, unexpired lease and current input digest, insert uniquely keyed completed items, persist cursor/progress and release lease. A retry may reuse verified item results but cannot publish a comparison until every required item has a supported disposition. Stale work is marked stale and does not update the current comparison.
- 6. Finalize assessment/agent packet:** verify dependency versions, result schema and host validation; write immutable result, set current pointer only if current, append audit and complete idempotency record in one transaction. A lease acquired before an input change cannot overwrite newer evidence.
- 7. Expire/reclaim:** an authorized resume request may mark expired work interrupted, release expired reservations according to policy and start a fresh lease. GET status never starts computation. Failed accounting uses conservative reservation until expiry; it must not allow double spending.
- 8. Reset workspace:** owner-only revision-protected reset creates a new seed generation and cancels/stales prior operations. Perform bounded deletion/archival with relational integrity; never reset another workspace or the whole database.

POST bodies do not carry executable SQL, object-storage destinations, unrestricted URLs or callback addresses. Runtime file access is limited to packaged assets and temporary bounded scratch data.

Queue time and acknowledgement details

Policy depends on the demo clock. Advance the scenario clock through a versioned workspace settings operation; increment queue revision and recompute affected priorities before serving the next queue snapshot. In live-wall-clock mode, detect the next deadline boundary on a request and atomically refresh stale policy buckets. A cursor is invalid when the resulting queue revision differs. Do not mutate the clock on every GET or silently reuse yesterday's urgency.

An urgent floor belongs to a particular unacknowledged material assessment chain. A downgrade proposal remains visible alongside the floor. An explicit latest-version acknowledgement records the actor/rationale and may clear the floor only if deterministic current policy permits the lower urgency. Closing a case does not cancel future material re-evaluation.

Migrations and security acceptance

Order migrations: enums/base tables; foreign keys/indexes; membership helpers/RLS/grants; domain RPCs; synthetic template seed; optional retention maintenance. Each migration is versioned and reviewed. Avoid destructive renaming/drops during a release; use additive fields and an explicit later migration if necessary.

Test through real Supabase Auth and Data API, not only mocked Python repositories: two identities, viewer role, forged actor, guessed UUID, cross-workspace foreign key, expired JWT, direct assessment insertion, direct RPC execution, simultaneous updates, lease expiry, interrupted response and idempotent retry. These are required because service-key or grant mistakes can make otherwise correct application checks ineffective.

Source notes

18. Supabase, [Row Level Security](#).